

# **Machine Learning** - Deep Learning fundamentals (69152) **CNNs**

---

*Master in Robotics, Graphics and Computer Vision*  
Ana C. Murillo



**Universidad**  
Zaragoza

# Today

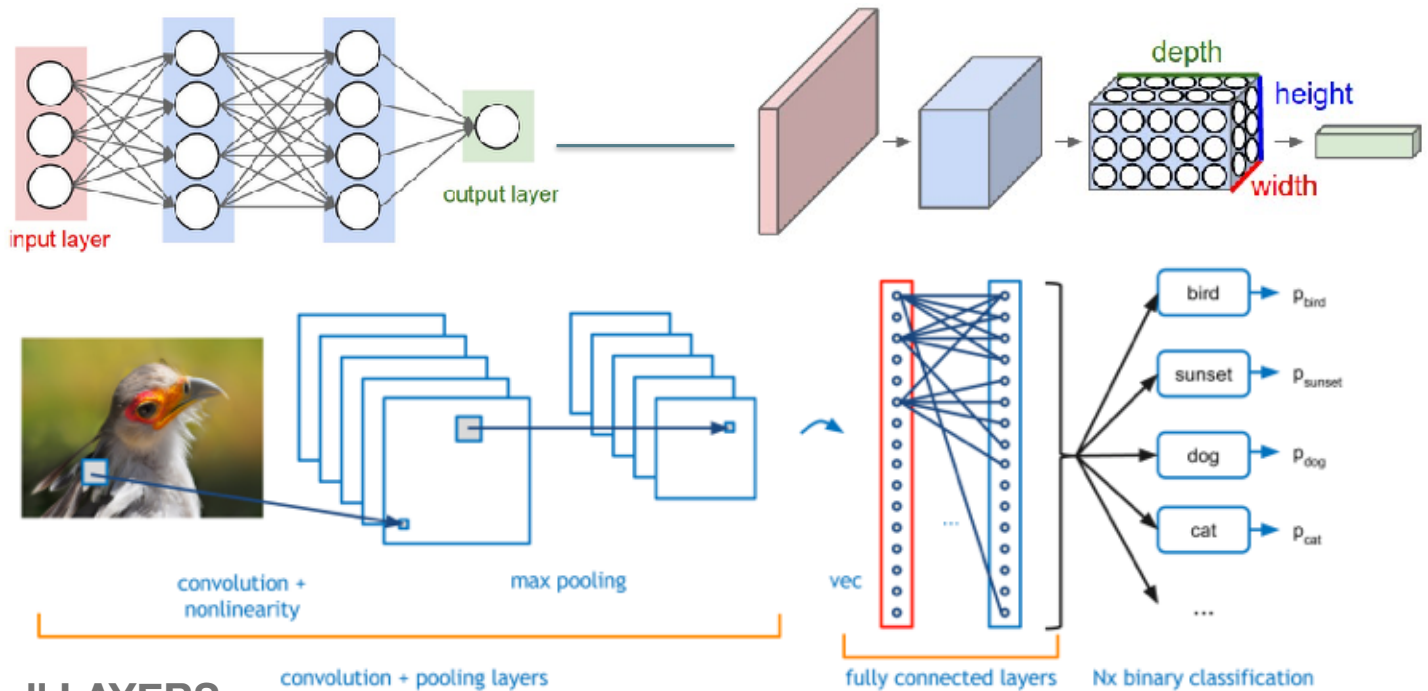
---

- What's a CNN?
- Well-known CNN architectures

# CNNs

- Convolutional (deep) Neural Networks (CNN)

[DEMO online, CNN](#)



**NOT all LAYERS  
are the same**


Fei-Fei, Karpathy, Johnson. Convolutional Neural Networks for Visual Recognition (<http://cs231n.stanford.edu>)  
Evan Shelhamer, Jeff Donahue, Jon Long, Yangqing Jia, and Ross Girshick. *Deep Learning for Vision: a Hands-On Tutorial*

# CNNs - Basic Layers

---


- **Input:** raw pixel values. RGB image 32x32, volume [32x32x3]
- **Convolutional:** compute output of neurons connected to local input regions (each computes dot product between their weights and the region). Larger volume [32x32x12].
- **RELU:** element-wise activation function (e.g. thresholding at zero). Volume unchanged ([32x32x12]).
- **Pooling:** downsampling operation. e.g. reduce volume to [16x16x12].
- **Fully-connected:** compute class scores. Each neuron in this layer will be connected to all the numbers in the previous volume. Volume of size [1x1x10].

# CNNs - pooling example



Fei-Fei, Karpathy, Johnson. Convolutional Neural Networks for Visual Recognition (<http://cs231n.stanford.edu>)  
Evan Shelhamer, Jeff Donahue, Jon Long, Yangqing Jia, and Ross Girshick. *Deep Learning for Vision: a Hands-On Tutorial*

# CNNs - pooling example



Fei-Fei, Karpathy, Johnson. Convolutional Neural Networks for Visual Recognition (<http://cs231n.stanford.edu>)  
Evan Shelhamer, Jeff Donahue, Jon Long, Yangqing Jia, and Ross Girshick. *Deep Learning for Vision: a Hands-On Tutorial*


# CNNs

---

- A bit of “history” about CNNs

# CNNs - Image Classification

**CNNs basics: AlexNet** (Beginning of the *new CNN boom*):  
*a layered model composed of convolution, subsampling, and further operations followed by a holistic representation.*



*Illustrated the  
benefits of CNNs*

- 15 million annotated images
- over 22000 classes
- **ReLU, data augmentation, dropout**

IMAGENET

*ImageNet Classification with Deep Convolutional Neural Networks*  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS **2012**




# CNNs

## CNNs *basics*: DeConvNet

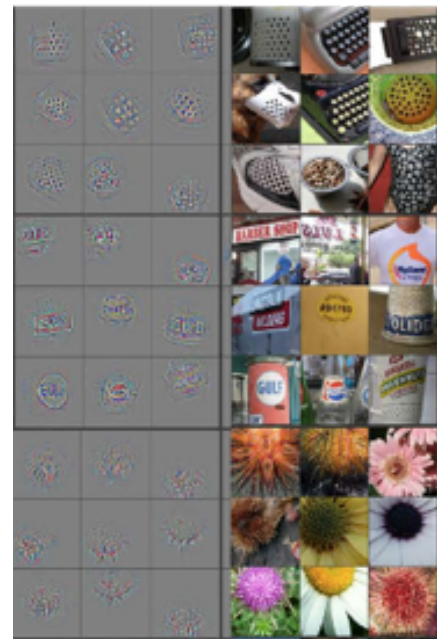


image patches that strongly activate 1st layer filters



1st layer filters

***better understanding:  
earlier vs later layers***



conv<sub>5</sub> DeConv visualization  
[Zeiler-Fergus]

- Similar architecture to AlexNet
- less training data - smaller filters
- Deconvolutional Network - **visualization**


*Visualizing and Understanding Convolutional Neural Networks*  
Matthew Zeiler and Rob Fergus. 2013

<http://people.csail.mit.edu/torralba/research/drawCNN/drawNet.html>

# CNNs

---

## CNNs *basics*: DeConvNet



1st layer filters

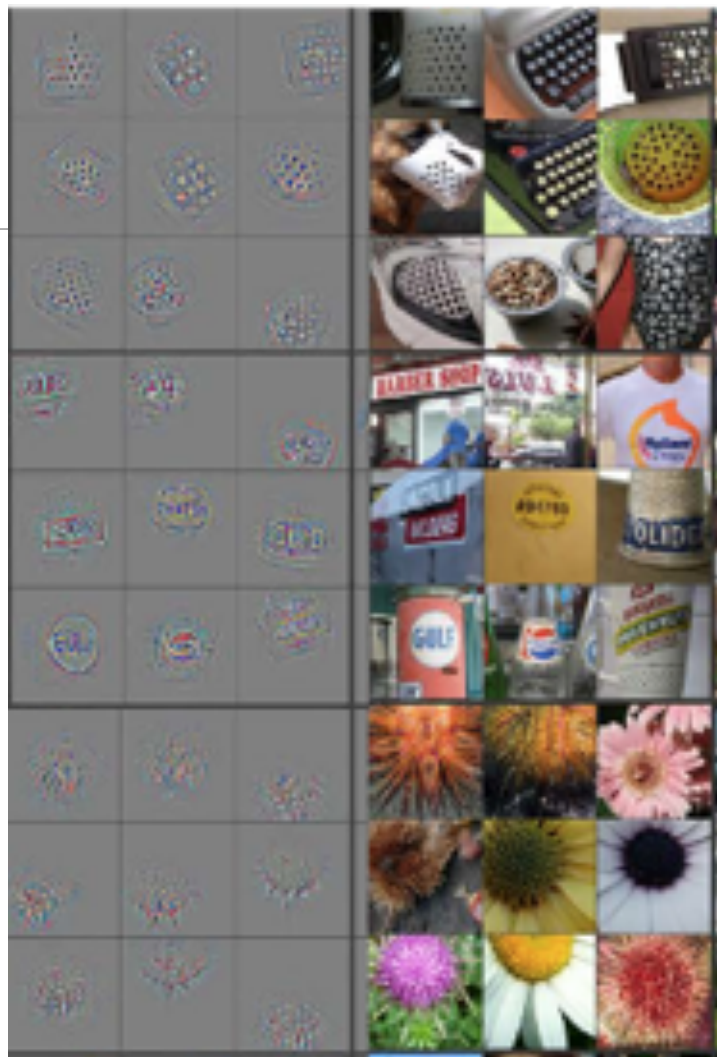
*Visualizing and Understanding Convolutional Neural Networks*  
Matthew Zeiler and Rob Fergus. **2013**



# CNNs

## CNNs *basics*: DeConvNet

***top 9 activations for a few feature maps***  
(projection to pixel space using DeConv  
and actual image patches)




conv<sub>5</sub> DeConv visualization  
[Zeiler-Fergus]

# CNNs - Image Classification + Detection

- **Analyze the *feature maps*: Region - CNN**

(R-CNN - 2013, Fast R-CNN - 2015, Faster R-CNN - 2016)



**Analyze feature maps  
(activations) to learn  
where the objects are**


*Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks*  
Shaoqing Ren, Kaiming He, **Ross Girshick**, and Jian Sun. 2016

# CNNs - Image classification

- *Deeper* classification: VGG - **very deep** ConvNets
  - Smaller filters - less params
  - more depth! —> always good (?)
  - Consecutive conv. layers (smaller filters).
  - Compare to AlexNet 5 conv. layers




How deep?



*Very Deep Convolutional Networks for Large-Scale Image Recognition.* K. Simonyan, A. Zisserman. **2014**  
*Deep Face Recognition:* O. M. Parkhi, A. Vedaldi, A. Zisserman. BMVC. **2015**

slide from R. Zemel. DLSS. 2017

# Image classification (deeper)



16 convolutional layers

50 convolutional layers

152 convolutional layers

Krizhevsky, I. Sutskever, and G. E. Hinton. *ImageNet classification with deep convolutional neural networks*. In Proc. NIPS, 2012.

C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich. *Going deeper with convolutions*. In Proc. CVPR, 2015.

K. Simonyan and A. Zisserman. *Very deep convolutional networks for large-scale image recognition*. In Proc. ICLR, 2015.

K. He, X. Zhang, S. Ren, and J. Sun. *Deep residual learning for image recognition*. In Proc. CVPR, 2016.

**Are they just deeper version of the same?**



# Deeper: How does memory get affected?

INPUT: [224x224x3] memory:  $224*224*3=150K$  params: 0 (not counting biases)

CONV3-64: [224x224x64] memory:  $224*224*64=3.2M$  params:  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory:  $224*224*64=3.2M$  params:  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory:  $112*112*64=800K$  params: 0

CONV3-128: [112x112x128] memory:  $112*112*128=1.6M$  params:  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory:  $112*112*128=1.6M$  params:  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory:  $56*56*128=400K$  params: 0

CONV3-256: [56x56x256] memory:  $56*56*256=800K$  params:  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory:  $56*56*256=800K$  params:  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory:  $56*56*256=800K$  params:  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory:  $28*28*256=200K$  params: 0

CONV3-512: [28x28x512] memory:  $28*28*512=400K$  params:  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory:  $28*28*512=400K$  params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory:  $28*28*512=400K$  params:  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory:  $14*14*512=100K$  params: 0

CONV3-512: [14x14x512] memory:  $14*14*512=100K$  params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory:  $14*14*512=100K$  params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory:  $14*14*512=100K$  params:  $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory:  $7*7*512=25K$  params: 0

FC: [1x1x4096] memory: 4096 params:  $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params:  $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params:  $4096*1000 = 4,096,000$

TOTAL memory:  $24M * 4 \text{ bytes} \sim 96MB$  / image (only forward!  $\sim 2$  for bwd)

TOTAL params: 138M parameters

Very Deep Convolutional Networks for Large-Scale Image Recognition.  
K. Simonyan, A. Zisserman. 2014

image
Conv-64
Conv-64
maxpool
Conv-128
Conv-128
maxpool
Conv-256
Conv-256
maxpool
Conv-512
Conv-512
Conv-512
maxpool
Conv-512
Conv-512
Conv-512
maxpool
fc-4096
fc-4096
fc-2622
Softmax

# Deeper: How does memory get affected?

image
Conv-64
Conv-64
maxpool
Conv-128
Conv-128
maxpool
Conv-256
Conv-256
maxpool
Conv-512
Conv-512
Conv-512
maxpool
Conv-512
Conv-512
Conv-512
maxpool
fc-4096
fc-4096
fc-2622
Softmax

INPUT: [224x224x3] memory: 224\*224\*3=150K params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: 224\*224\*64=3.2M params: (3\*3\*3)\*64 = 1,728

CONV3-64: [224x224x64] memory: 224\*224\*64=3.2M params: (3\*3\*64)\*64 = 36,864

POOL2: [112x112x64] memory: 112\*112\*64=800K params: 0

CONV3-128: [112x112x128] memory: 112\*112\*128=1.6M params: (3\*3\*64)\*128 = 73,728

CONV3-128: [112x112x128] memory: 112\*112\*128=1.6M params: (3\*3\*128)\*128 = 147,456

POOL2: [56x56x128] memory: 56\*56\*128=400K params: 0

CONV3-256: [56x56x256] memory: 56\*56\*256=800K params: (3\*3\*128)\*256 = 294,912

CONV3-256: [56x56x256] memory: 56\*56\*256=800K params: (3\*3\*256)\*256 = 589,824

CONV3-256: [56x56x256] memory: 56\*56\*256=800K params: (3\*3\*256)\*256 = 589,824

POOL2: [28x28x256] memory: 28\*28\*256=200K params: 0

CONV3-512: [28x28x512] memory: 28\*28\*512=400K params: (3\*3\*256)\*512 = 1,179,648

CONV3-512: [28x28x512] memory: 28\*28\*512=400K params: (3\*3\*512)\*512 = 2,359,296

CONV3-512: [28x28x512] memory: 28\*28\*512=400K params: (3\*3\*512)\*512 = 2,359,296

POOL2: [14x14x512] memory: 14\*14\*512=100K params: 0

CONV3-512: [14x14x512] memory: 14\*14\*512=100K params: (3\*3\*512)\*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14\*14\*512=100K params: (3\*3\*512)\*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14\*14\*512=100K params: (3\*3\*512)\*512 = 2,359,296

POOL2: [7x7x512] memory: 7\*7\*512=25K params: 0

FC: [1x1x4096] memory: 4096 params: 7\*7\*512\*4096 = 102,760,448

FC: [1x1x4096] memory: 4096 params: 4096\*4096 = 16,777,216

FC: [1x1x1000] memory: 1000 params: 4096\*1000 = 4,096,000

Note:

Most memory is in early CONV

Most params are in late FC

TOTAL memory: 24M \* 4 bytes ~= 96MB / image (only forward! ~\*2 for bwd)

TOTAL params: 138M parameters

*What is more relevant for batch size design decisions?*


Very Deep Convolutional Networks for Large-Scale Image Recognition.  
K. Simonyan, A. Zisserman. 2014

slide from  
Fei-Fei Li & Justin Johnson & Serena Yeung  
<http://cs231n.stanford.edu>




# Image classification (deeper)

- Not only more layers! **it's getting hard(er) to train ...**



ImageNet Classification with Deep Convolutional Neural Networks  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012



16 convolutional layers →  
50 convolutional layers →  
152 convolutional layers →

Krizhevsky, I. Sutskever, and G. E. Hinton. *ImageNet classification with deep convolutional neural networks*. In Proc. NIPS, 2012.

C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich. *Going deeper with convolutions*. In Proc. CVPR, 2015.


K. Simonyan and A. Zisserman. *Very deep convolutional networks for large-scale image recognition*. In Proc. ICLR, 2015.

K. He, X. Zhang, S. Ren, and J. Sun. *Deep residual learning for image recognition*. In Proc. CVPR, 2016.

slide from R. Zemel. DLSS. 2017

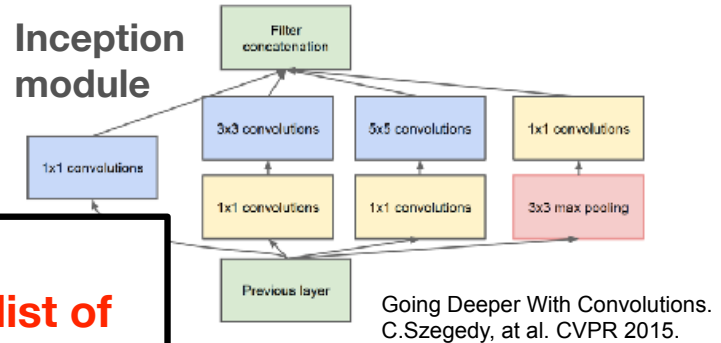
# Some *reference modules* for CNN architectures ...

- to train better: GoogLeNet, ResNet, DenseNet, ...




ImageNet Classification with Deep Convolutional Neural Networks  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2015

## Inception module




Going Deeper With Convolutions.  
C.Szegedy, et al. CVPR 2015.

## Residual module




Deep residual learning for image recognition.  
He, Kaiming, et al. CVPR 2016



Not just  
"linear" list of  
layers anymore

## Dense module



Densely Connected Convolutional Networks.  
Huang, G., et al. CVPR 2017.

Open or interesting problems related to deep learning?

**Not enough data to train? (or resources)**


# CNNs & Transfer Learning

- CNNs are able to generalize well!

- great **features**

- **fine-tuning**

Lab 2  
you'll practice  
some of this



**deep features**

ImageNet Classification with Deep Convolutional Neural Networks  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012


# Transfer Learning

---


- What is it?
- When would you use transfer learning?
- Types of transfer learning?

# Transfer Learning: features

- E.g., features from AlexNet last layer: 4096 dims. vector



Test image L2 Nearest neighbors in feature space

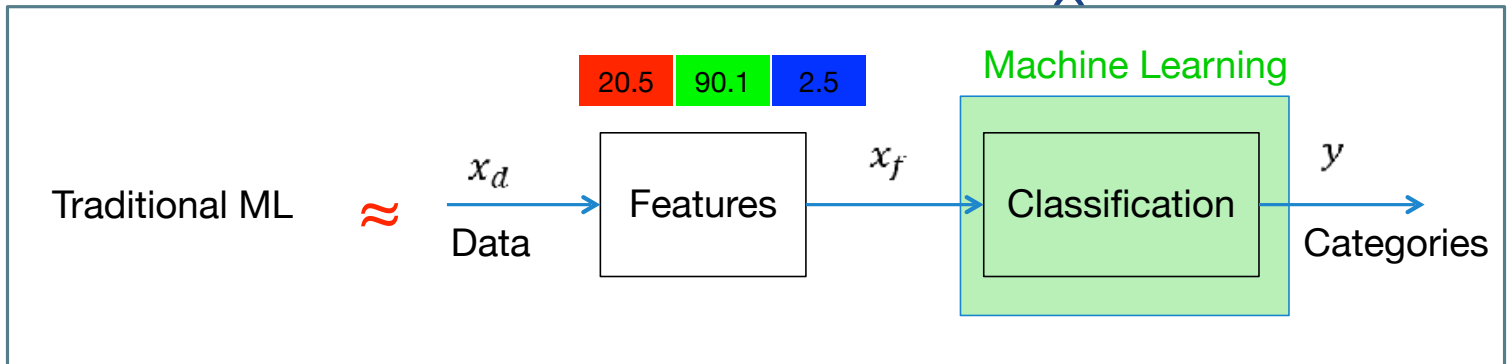


***How do we get these features?***

ImageNet Classification with Deep Convolutional Neural Networks  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012


# Transfer Learning: features

- E.g., features from AlexNet last layer: 4096 dims. vector



***How do we use these features?***


Test image L2 Nearest neighbors in feature space



ImageNet Classification with Deep Convolutional Neural Networks  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012

# Transfer Learning: fine-tune

- For example, fine-tune ImageNet AlexNet for non Image-Net classes
- Basically, initialise weights to something more “*interesting*” than random (careful also with hyperparameters!)

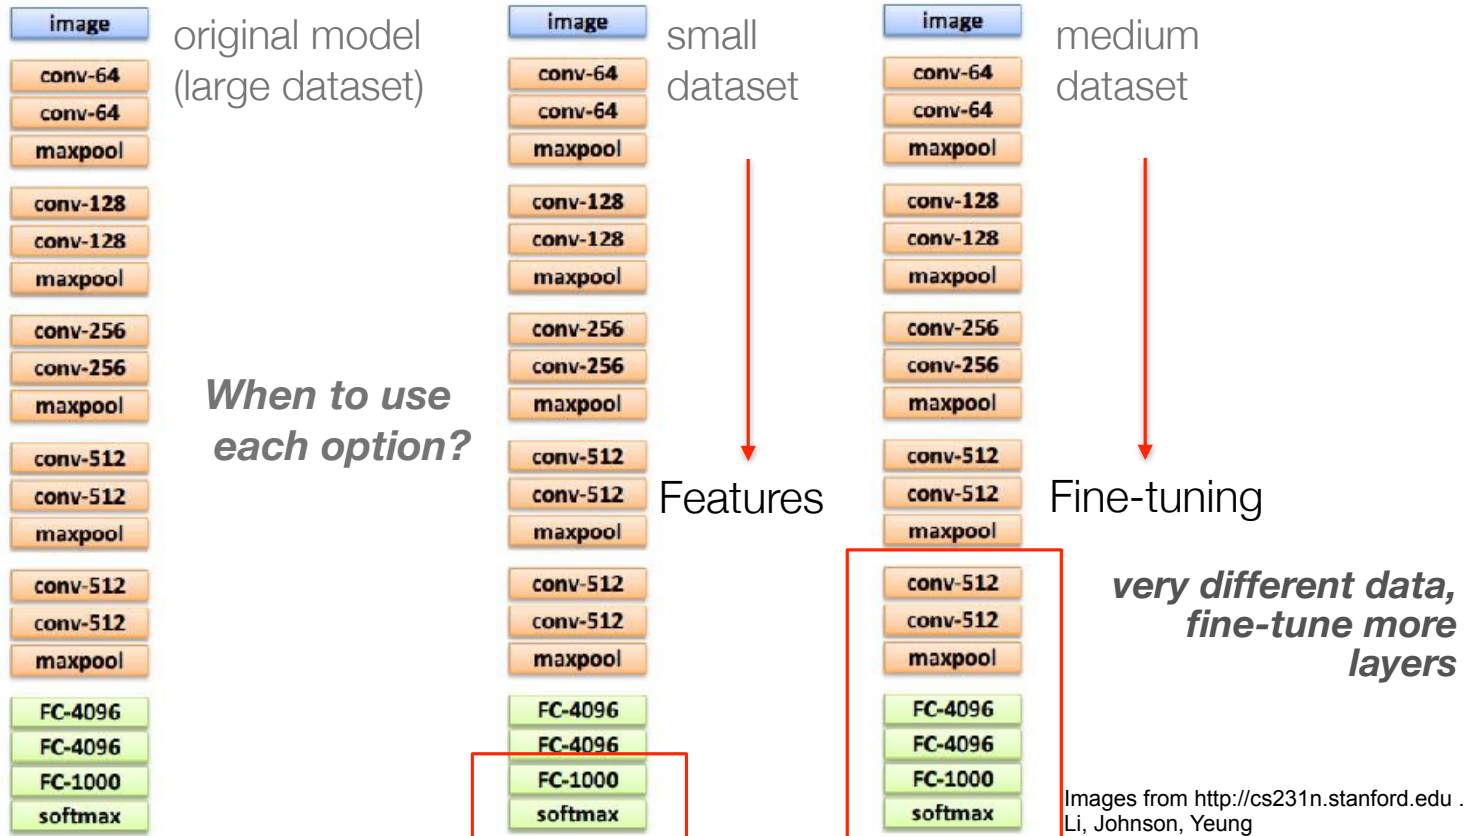


ImageNet Classification with Deep Convolutional Neural Networks  
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012



# Transfer Learning

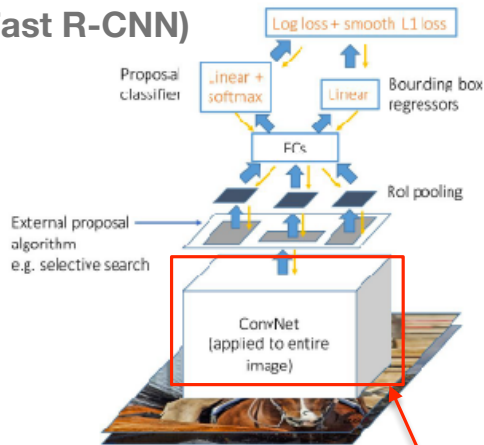
... when not enough resources to train from scratch



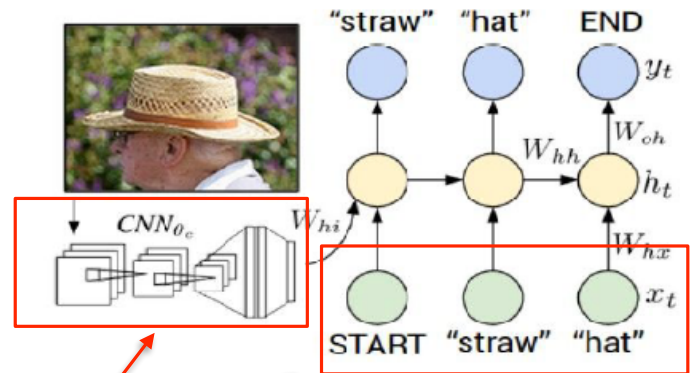
# Transfer Learning: fine-tuning

- not just that! very very spread out strategy:

## Object Detection (Fast R-CNN)



## Image captioning: CNN + RNN



*Pre-trained  
models to  
initialize certain  
blocks*

Karpathy and Fei-Fei, "Deep Visual-Semantic Alignments for Generating Image Descriptions", CVPR 2015  
Figure copyright IEEE, 2015. Reproduced for educational purposes.

Images from <http://cs231n.stanford.edu> . Li, Johnson, Yeung

Girshick, "Fast R-CNN", ICCV 2015  
Figure copyright Ross Girshick, 2015. Reproduced with permission.

# Lab 2: NN and CNNs applied in Computer Vision

---

- Understand a **DNN implemented from scratch**
- Implement a toy CNN with Keras
- Fine-tune a well-known CNN architecture with Keras

All this is applied to Image Classification

- Optional tasks:
  - implement variations of the classification networks
  - explore object recognition model
  - explore semantic segmentation model

***Available in Moodle.***

***Please check the prior-work task (prepare your set up and data)***

## TO-DO ...

*Have you all used COLAB?*

- Lab 2 - **TOMORROW (18 OCT , 20 OCT) —> WED. @ A07 / FRIDAY @ L0.06a ADA BYRON**
- PLEASE **have your computer ready with Tensorflow2+Keras (ideally with GPU available) AND/OR we will use Google COLAB**
- Recommended COLAB tutorial if you have not used it much before:  
<https://colab.research.google.com/notebooks/intro.ipynb>  
[https://colab.research.google.com/notebooks/basic\\_features\\_overview.ipynb](https://colab.research.google.com/notebooks/basic_features_overview.ipynb)  
[Loading data: Drive, Sheets, and Google Cloud Storage](#)

## ASSIGNMENT BEFORE YOUR LAB:

- Pick 5 to 10 classes from one of these datasets (do not take all images from each class if you don't have space)
  - <https://www.kaggle.com/kmader/food41/version/5#>
  - <http://www.robots.ox.ac.uk/~vgg/data/pets/>
  - <http://www.robots.ox.ac.uk/~vgg/data/flowers/>
  - any other dataset you have?**
- Put them in folders like —————>
- Upload to Google Drive if you plan to use COLAB

```
data/  
  dogs/  
    dog001.jpg  
    dog002.jpg  
    ...  
  cats/  
    cat001.jpg  
    cat002.jpg  
    ...
```

## Lab 2: NN and CNNs applied in Computer Vision

---

- Frameworks
  - Caffe, **Tensorflow+Keras**, **Pytorch**, ...
- Model **zoos**
  - in every framework (*you should always start by looking into existing models. “rare” to “require” to implement a network from scratch*)

# Understanding the training process


---

Initial sanity checks (after a few “toy” iterations)


- *Correct* loss at chance
- Increasing regularisation strength increases loss
- Overfit a tiny training set (you can set regularisation strength to zero to better see this)

# Understanding the training process

- Monitor the loss function




Typical loss function.  
Batch size a little low? (noisy plot)



# Understanding the training process

---

- Monitor train/val accuracy







# Understanding the training process

---

- Visualize first conv layer weights



# Understanding the training process

---

Additional tips:

- **Decay your learning rate** during training
- If you can afford it
  - **Search** for good **hyperparameters** with random search (not grid search). From coarse to fine
  - Consider **model ensembles** (not in our labs)

# Lab 2: NN and CNNs applied in Computer Vision

---

- Understand a **DNN implemented from scratch**
- Implement a toy CNN with Keras
- Fine-tune a well-known CNN architecture with Keras

All this is applied to Image Classification

- Optional tasks:
  - implement variations of the classification networks
  - explore object recognition model
  - explore semantic segmentation model

***Exercise now  
(start for Lab2)***

**[COLAB](#)**