

# Reinforcement Learning

## Machine Learning (69152)

**Rubén Martínez Cantín**

Dpto. Informática e Ingeniería de Sistemas  
Universidad de Zaragoza

## 1 Introduction

- Markov Decision Processes (recap)
- Learning by interaction


## 2 Prediction

- Monte Carlo Evaluation
- Temporal Differences

## 3 Control

- Sampling control
- Exploration vs Exploitation
- Q-learning
- Q-learning with features

# Markov Decision Processes




In particular, for **Markov Decision Processes (MDPs)**:

- Markov property  $x_{t+1} = f(x_t)$
- Full observability  $y_t = x_{t+1}$
- Action  $a_t \Leftarrow \pi(x_t)$

# Markov Decision Process

- A Markov Decision Process is a tuple  $\langle \mathcal{X}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$

- ▶ A set of states  $x \in \mathcal{X}$ .
- ▶ A set of actions  $a \in \mathcal{A}$ .
- ▶ A transition function  $T(x, a, x') = p(x'|x, a)$
- ▶ A reward function  $R(x, a, x')$
- ▶ Maybe a start state and a terminal state.



Credit: Dan Klein, Pieter Abbeel

# Markov Decision Process

- A Markov Decision Process is a tuple  $\langle \mathcal{X}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$

- ▶ A set of states  $x \in \mathcal{X}$ .
- ▶ A set of actions  $a \in \mathcal{A}$ .
- ▶ A transition function
$$T(x, a, x') = p(x'|x, a)$$
- ▶ A reward function  $R(x, a, x')$
- ▶ Maybe a start state and a terminal state.



- **Plot twist:** Now we don't know the transitions  $T(\cdot)$  or rewards  $R(\cdot)$ 
  - ▶ The agent must try actions to see the outcome.

Credit: Dan Klein, Pieter Abbeel

# Learning by interaction

## Model-based reinforcement learning

- Learn an approximate model based on experiences
- Assume that the learned model is correct and solve it as if known MDP.

## Model-free reinforcement learning

- Learn the policy and values directly from interaction
- No explicit model of the transitions  $T(\cdot)$  or rewards  $R(\cdot)$

## Example: Compute the expected age

We want to compute the expected age of all the students in the University.

Known Model  $p(a)$

$$\mathbb{E}(A) = \sum_a p(a) * a$$

If we  $p(a)$  is unknown we can ask several students for the age  $[a_1, a_2, \dots]$

Model based

$$\hat{p}(a) = \frac{\text{num}(a)}{N}$$
$$\mathbb{E}(A) \approx \sum_a \hat{p}(a) * a$$

Model free

$$\mathbb{E}(A) \approx \frac{\sum_i a_i}{N}$$

# Model based learning

- Model-Based Idea:
  - ▶ Learn an approximate model based on experiences
  - ▶ Solve for values as if the learned model were correct
- Step 1: Learn empirical MDP model
  - ▶ Count outcomes  $x'$  for each  $x, a$
  - ▶ Normalize to give an estimate of  $\hat{p}(x'|x, a)$
  - ▶ Discover each  $\hat{R}(x, a, x')$  when we experience  $(x, a, x')$
- Step 2: Solve the learned MDP
  - ▶ For example, use value iteration, as before



Credit: Dan Klein, Pieter Abbeel



# Model-free learning

- **Model-free Idea:** why bother learning the transition or reward model?
  - ▶ Can we compute  $V$ ,  $Q$  or  $\pi^*$  without a model?
  - ▶ **Monte Carlo:** Sample episodes, average rewards at the end.
  - ▶ **Temporal differences:** Use sampling to approximate the Bellman updates, compute new values during each learning step.



Credit: Dan Klein, Pieter Abbeel

# Model-free learning

- Passive reinforcement learning
  - ▶ Evaluation given a policy  $\pi$ , find the value  $V^\pi$ .
  - ▶ *Learner* has no choice. Just follow the policy
  - ▶ The agent is doing **prediction**.
- Active reinforcement learning
  - ▶ Find the optimal policy/values:  $V^*, Q^*, \pi^*$ .
  - ▶ *Learner* can choose. Exploration vs Exploitation.
  - ▶ The agent is doing **control**.

# Model-free learning

- Passive reinforcement learning
  - ▶ Evaluation given a policy  $\pi$ , find the value  $V^\pi$ .
  - ▶ *Learner* has no choice. Just follow the policy
  - ▶ The agent is doing **prediction**.
- Active reinforcement learning
  - ▶ Find the optimal policy/values:  $V^*, Q^*, \pi^*$ .
  - ▶ *Learner* can choose. Exploration vs Exploitation.
  - ▶ The agent is doing **control**.
- Remember: even if the learner can choose or not, it must take actions and interact with the world.

# Monte Carlo Direct Evaluation

- Monte Carlo (MC) passive reinforcement learning.
- Goal: learn  $V^\pi$  from episodes of experience under policy  $\pi$
- Act according to  $\pi$ .

$$x_1, a_1, R_2, x_2, \dots, x_n \sim \pi$$

- Recall that the utility is the total discounted reward:

$$U_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots R_n$$

- Recall that the value function is the expected return:

$$V^\pi(x) = \mathbb{E}_\pi[U_t | x_t = x]$$

- MC policy evaluation uses empirical mean return instead of expected return


$$V^\pi(x) = \sum_i U_t^{(i)} \mathbb{1}_{x_t^{(i)}=x} \quad \forall i \in \text{episodes}$$

# Monte Carlo Direct Evaluation

- To evaluate state  $x$
- Every time-step  $t$  that state  $x$  is visited in an episode,
- Increment counter  $N(x) \leftarrow N(x) + 1$
- Increment total return  $S(x) \leftarrow S(x) + U_t$
- Value is estimated by mean return  $\hat{V}(x) = S(x)/N(x)$
- By law of large numbers,  $\hat{V}(x) \rightarrow V^\pi(x)$  as  $N(x) \rightarrow \infty$

# Monte Carlo Direct Evaluation

Input policy  $\pi$



Assume:  $\gamma = 1$

Observed Episodes

Episode 1

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 3

- E, north, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 2

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 4

- E, north, C, -1
- C, east, A, -1
- A, exit, x, -10


Estimated values

|    |     |     |
|----|-----|-----|
|    | -10 |     |
| +8 | +4  | +10 |
|    | -2  |     |

Credit: Dan Klein, Pieter Abbeel

# Monte Carlo Direct Evaluation

Input policy  $\pi$



Assume:  $\gamma = 1$

Observed Episodes

Episode 1

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 3

- E, north, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 2

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 4

- E, north, C, -1
- C, east, A, -1
- A, exit, x, -10

Estimated values

|    |     |     |
|----|-----|-----|
|    | -10 |     |
| +8 | +4  | +10 |
|    | -2  |     |

- How can B and E have different values if both go to C for this policy?

Credit: Dan Klein, Pieter Abbeel

# Incremental Monte Carlo

- Do we need to wait until the end of all episodes to compute the empirical mean?
- Remember the *expected age* example:
  - ▶ Model free approach:  $\mathbb{E}(A) \approx \hat{A}_N = \frac{\sum_{i=1}^N a_i}{N}$
  - ▶ Incremental approach:

$$\begin{aligned}\hat{A}_N &\approx \frac{1}{N} \sum_{i=1}^N a_i \\ &= \frac{1}{N} \left( a_N + \sum_{i=1}^{N-1} a_i \right) \\ &= \frac{1}{N} \left( a_N + (N-1)\hat{A}_{N-1} \right) \\ &= \hat{A}_{N-1} + \frac{1}{N} \left( a_N - \hat{A}_{N-1} \right)\end{aligned}$$



# Incremental Monte Carlo for value function

- Update  $V(x)$  incrementally after episode  $x_1, a_1, R_2, x_2, \dots, x_n$
- For each state  $x_t$  with utility  $U_t$

$$N(x_t) \leftarrow N(x_t) + 1$$

$$V(x_t) \leftarrow V(x_t) + \frac{1}{N(x_t)}(U_t - V(x_t))$$

- In non-stationary problems, it can be useful to track a running mean, that is, forget old episodes.

$$V(x_t) \leftarrow V(x_t) + \alpha(U_t - V(x_t))$$

# Problems with Monte Carlo Direct Evaluation

- Samples are based on full episodes.
- The problem needs to be episodic. Not for continuous.
- It ignores MDP structure.
- Is there a way to exploit MDP structure and use Bellman equations?

# Monte Carlo Policy Evaluation

- We had the Bellman expectation equation to update the value function:

$$V^{\pi}(x) = \sum_{a \in \mathcal{A}} \pi(a|x) \left( \sum_{x' \in \mathcal{X}} p(x'|x, a) (R(x, a, x') + \gamma V^{\pi}(x')) \right)$$

- This approach exploits the MDP structure and is valid for continuous problems. 😊
- But it requires  $p(x'|x, a)$  and  $R(x, a, x')$  to do it. 😞

# Sample based policy evaluation

- Instead of sampling full episodes, we can sample the value update.

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(a|x) \left( \sum_{x' \in \mathcal{X}} p(x'|x, a) (R(x, a, x') + \gamma V^\pi(x')) \right)$$

- From one state  $x$ , we sample just the next action  $a_1, a_2, \dots \sim \pi(x)$  and the next transition  $x_1 \sim p(x'|a_1, x), x_2 \sim p(x'|a_2, x), \dots$
- Then, we collect the outcome for the action/transition:

$$sample_1 = R(x, a_1, x'_1) + \gamma V_k^\pi(x'_1)$$

$$sample_2 = R(x, a_2, x'_2) + \gamma V_k^\pi(x'_2)$$

...

- Finally,

$$V_{k+1}^\pi(x) \leftarrow \frac{1}{N} \sum_{i=1}^N sample_i$$

# Temporal Difference

- **Idea:** learn from any interaction.
- Combine the sample based policy evaluation with the incremental Monte Carlo method.
- Remember: Incremental Monte Carlo updates the value with the *actual* utility  $U_t$ .

$$V(x_t) \leftarrow V(x_t) + \alpha(U_t - V(x_t))$$

- Temporal Difference updates the value with the *estimated* utility  $R_{t+1} + \gamma V(x_{t+1})$

$$V(x_t) \leftarrow V(x_t) + \alpha(R_{t+1} + \gamma V(x_{t+1}) - V(x_t))$$

# Temporal Difference

- Temporal Difference (TD) can be learned after any interaction. No need to wait for the end of the episode.
- It can be used for continuous problems.
- It can even learn from incomplete sequences.
- It exploits MDP structure.

$$\begin{aligned} \text{sample}_k &= R(x, a, x') + \gamma V_k^\pi(x') \\ V_{k+1}^\pi(x) &\leftarrow (1 - \alpha) V_k^\pi(x) + \alpha \cdot \text{sample}_k = \\ &= V_k^\pi(x) + \alpha (\text{sample}_k - V_k^\pi(x)) \end{aligned}$$

- Remember that  $\alpha \in [0, 1]$  is used to *forget*.
  - $\gamma \rightarrow 0$  ignores the distant future.
  - $\alpha \rightarrow 1$  forgets the distant past

# Forgetting the past

- Remember, we can use a running mean to forget old updates.
  - ▶ In MC, for non-stationary problems
  - ▶ In TD, recursive update of Bellman equation:

$$V(x_t) \leftarrow V(x_t) + \alpha(\text{update}_t - V(x_t))$$

- A running average  $\hat{z}$  of  $\{z\}_{i=1}^N$  elements:

$$\hat{z}_k = \hat{z}_{k-1} + \alpha(z_k - \hat{z}_{k-1}) = (1 - \alpha)\hat{z}_{k-1} + \alpha z_k$$

- Weights recent samples more:

$$\hat{z}_k = \frac{z_k + (1 - \alpha)z_{k-1} + (1 - \alpha)^2 z_{k-2} + \dots}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots}$$


# MC vs TD

- Both methods converge to the true value if experience  $\rightarrow \infty$
- MC requires full episodes.
- MC is a very simple and general idea. It can be applied for any system.
- TD exploits the Markov property.
  - ▶ It is more efficient for Markov environments.
  - ▶ It might not converge for certain environments (e.g.: non-Markov).



# Monte Carlo Update


$$V(x_t) \leftarrow V(x_t) + \alpha(U_t - V(x_t))$$



Credit: David Silver

# Temporal Difference Update


$$V(x_t) \leftarrow V(x_t) + \alpha(R_{t+1} + \gamma V(x_{t+1}) - V(x_t))$$



Credit: David Silver

# Dynamic Programming Update

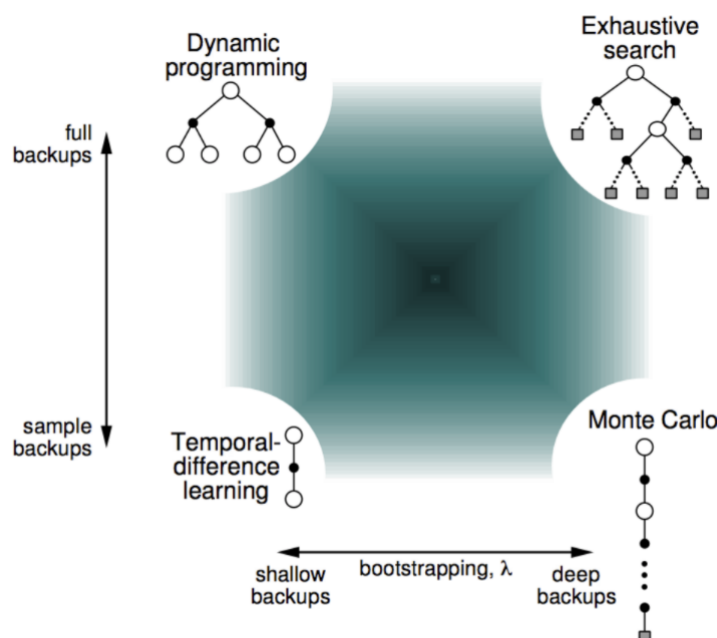
$$V(x_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(x_{t+1})]$$



Credit: David Silver

# Unified view of RL

- The TD algorithm that we have seen is called  $TD(0)$ , because it considers only one step ahead.
- The generalization is called  $TD(\lambda)$ , which combines multiple steps ahead.



Credit: David Silver

# Sampling control

- We have seen how to use MC and TD for **prediction** (policy evaluation).
- We are going to use MC and TD for **control** (optimal policy/value).
- **On-policy** learning
  - ▶ *Learn on the job*
  - ▶ Learn about policy  $\pi$  from experience sampled from  $\pi$
- **Off-policy** learning
  - ▶ *Look over someone's shoulder*
  - ▶ Learn about policy  $\pi$  from experience sampled from  $\mu$

## Quick Parenthesis: Q-values

- Computing the greedy policy from  $V(x)$  requires the MDP model  $p(x'|x, a)$  and  $R(x, a, x')$ .

$$\pi'(x) = \arg \max_a \sum_{x'} p(x'|x, a) (R(x, a, x') + \gamma V(x'))$$



- Computing the greedy policy from  $Q(x, a)$  can be done model-free:

$$\pi'(x) = \arg \max_a Q(x, a)$$

- **Solution:** we are going to do control based on Q-functions.


# SARSA

- Iteratively estimate  $Q$  and  $\pi$  (similar to policy iteration)
  - Policy evaluation (estimate  $Q$  given  $\pi$ ):
    - Good: Monte Carlo: run multiple full episodes, then update policy.
    - Better: Improve policy **after each episode**.
    - Best: Do not run full episodes. Use Temporal Differences.
- $$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R + \gamma Q(x_{t+1}, a_{t+1}) - Q(x_t, a_t))$$
- Policy improvement (improve  $\pi$  given  $Q$ ):
    - Exploit: The optimal policy is **greedy** to the value function...
    - Explore: ... but remember that we also need to **explore**!



# Exploration vs Exploitation

- The greedy policy is the best option if we know the model.
- Now, we are learning by experience.
  - ▶ The agent must have diverse experiences to learn.
  - ▶ But the agent must also find an optimal behavior.





# Regret

- Regret measures your **mistakes**. The (expected) reward of your actions, including suboptimal choices, against the (expected) optimal reward.
- Minimizing regret is not only learning to be optimal. It is **optimally learning to be optimal**.
- Pure random exploration will always find the optimal policy, but it has very high regret.
- A good trade-off of exploration and exploitation can provide **optimal regret**.

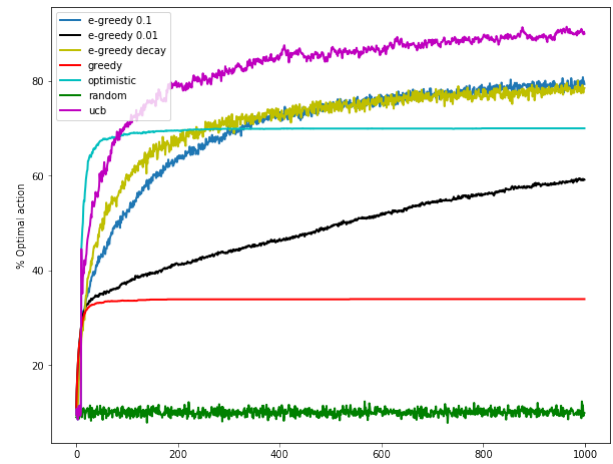




# Exploration-exploitation trade-off principles

- The limit cases are the *greedy policy* (always exploit) and the *random policy* (always exploit).
  - ▶ We have seen that both policies have high regret.
- Random mixing ( $\epsilon$ -greedy)
  - ▶ The simplest approach: flip a coin and choose *greedy* or *random policy* depending on the outcome.
- Optimistic initialization
  - ▶ Assume *unknown = best*. Try everything at least once.
- Optimism in the Face of Uncertainty
  - ▶ In case of uncertainty, assume the best possible outcome.
  - ▶ For example, upper confidence bound
  - ▶ If you are right: you win!
  - ▶ If you are wrong: you learn a lot!

# Exploration-exploitation demo

Exercise/demo: <https://drive.google.com/file/d/1xAph2c1P8pPcJudiWBCd9Rrqt4p2E0Ra/view?usp=sharing>



# DP vs TD algorithms

|                                  | Full backup       | TD Backup   |
|----------------------------------|-------------------|-------------|
| $V^\pi(x), Q^\pi(x, a)$          | Policy evaluation | TD learning |
| $\pi^*(x, a), V^*(x), Q^*(x, a)$ | Policy iteration  | SARSA       |
| $V^*(x), Q^*(x, a)$              | Value iteration   | ???         |

# DP vs TD algorithms

|                                  | Full backup       | TD Backup   |
|----------------------------------|-------------------|-------------|
| $V^\pi(x), Q^\pi(x, a)$          | Policy evaluation | TD learning |
| $\pi^*(x, a), V^*(x), Q^*(x, a)$ | Policy iteration  | SARSA       |
| $V^*(x), Q^*(x, a)$              | Value iteration   | Q-learning  |

# Q-learning

- Start with Q-value iteration

$$Q_{k+1}(x_t, a_t) = \sum_{x_{t+1}} p(x_{t+1}|x_t, a_t) \left( R(x_t, a_t, x_{t+1}) + \gamma \max_{a'} Q_k(x_{t+1}, a') \right)$$

- Consider the sample update

$$sample = R(x_t, a_t, x_{t+1}) + \gamma \max_{a'} Q_k(x_{t+1}, a')$$

- Which, using the running average, can be incorporated as:

$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha (R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t))$$

# Q-learning algorithm

---

## Algorithm 1 Q-learning algorithm

---

**Input:** Step size  $\alpha \in [0, 1]$ , policy parameters: e.g., small  $\epsilon > 0$ .

Initialize  $Q(x, a)$ , for all  $x \in \mathcal{X}$ ,  $a \in \mathcal{A}$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

**for** each episode **do**:

    Initialize  $x_0$

**for** each step  $t$  in episode **do**:

        Choose  $a_t$  from  $x_t$  using a policy derived from  $Q$  (e.g.:  $\epsilon$ -greedy)

        Take action  $a_t$ , observe  $R_{t+1}$  and  $x_{t+1}$ .

$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t))$

**if**  $x_{t+1}$  is terminal **then** Stop episode

---

# On-policy vs Off-policy

- SARSA and Q-learning require 2 actions for each update.

$$Q(x_t, \mathbf{a}_t) \leftarrow Q(x_t, \mathbf{a}_t) + \alpha(R_{t+1} + \gamma Q(x_{t+1}, \mathbf{a}') - Q(x_t, \mathbf{a}_t))$$

- This may come from different sources:
  - ▶ The **behavior policy** is how the agent is acting.
  - ▶ The **target policy** is the policy that the agent is learning.
- On-policy vs off-policy
  - ▶ On-policy methods use the same policy for the behavior and target.
  - ▶ Off-policy methods use a different policy. They can learn the optimal policy even acting suboptimally!



# On-policy vs Off-policy

- In SARSA, all actions are selected according to the target policy  $\pi$ , for example:  $\epsilon$ -greedy.


$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R_{t+1} + \gamma Q(x_{t+1}, a_{t+1}) - Q(x_t, a_t))$$

- In Q-learning, actions are selected according to the **behavior policy**  $\mu$ , for example:  $\epsilon$ -greedy, but the agent assumes that in the future it will be optimal ( $\pi$  is the greedy policy).

$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t))$$


# Scaling Q-learning

- How many states does smallGrid Pacman have?



# Scaling Q-learning

- How many states does smallGrid Pacman have?




- Classic Pacman has:  $240^5 \cdot 2^{240} = 1.766847110^{72}$  states
- Basic Q-learning keeps a table with a Q-value for every state and action.
- This idea is not able to scale:
  - ▶ During training, it is impossible to visit all the states/actions.
  - ▶ The table cannot be kept in memory.


# Generalization Q-learning

- We want to learn from few states and generalize to **similar states**.


We learn that this is a  
*bad state*



Basic Q-learning tell  
us nothing about this  
state.



Or even this state!




- What do they have in common?

Credit: Dan Klein, Pieter Abbeel

# Approximate Q-learning

- Idea: exploit certain properties (features) of state.
- A feature is a function from states or (state,action) to numbers.
- We can use features to describe other functions
- Examples of features:
  - ▶ Distance to a ghost.
  - ▶ Distance to a dot.
  - ▶ Number of ghost.
  - ▶ Is Pacman in a corner/tunnel?
  - ▶  $1/(distance\_to\_dot)^2$
  - ▶ ...
  - ▶ Is it the state in this slide?
  - ▶ ...
  - ▶ Does this action get me closer to food?



Credit: Dan Klein, Pieter Abbeel

# Linear function approximation

- Using a feature representation, we can write a q function (or value function) for any state using a few weights:

$$V(x) = w_1 f_1(x) + w_2 f_2(x) + \dots + w_n f_n(x)$$
$$Q(x, a) = w'_1 f_1(x, a) + w'_2 f_2(x, a) + \dots + w'_n f_n(x, a)$$

- Advantage: our experience is summed up in a few powerful numbers
- Disadvantage: states may share features but actually be very different in value!

# Approximate Q-learning

- Take a linear Q-value function:

$$Q(x, a) = w_1 f_1(x, a) + w_2 f_2(x, a) + \dots + w_n f_n(x, a)$$

- Q-learning update is based on estimation error:

$$\text{error} = \left( R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') \right) - Q(x, a)$$

- Basic Q-learning updates each value with the error:

$$Q(x, a) \leftarrow Q(x, a) + \alpha \cdot \text{error}$$

- Approximate Q-learning updates the weights of active features:

$$w_i \leftarrow w_i + \alpha \cdot \text{error} \cdot f_i(x, a)$$

# Approximate Q-learning

- Approximate Q-learning updates the weights of active features:

$$w_i \leftarrow w_i + \alpha \cdot \text{error} \cdot f_i(x, a)$$

- Intuitive interpretation:
  - ▶ If something *unexpectedly* bad happens, penalize the features that were on  $\Rightarrow w_i \downarrow$ .
  - ▶ If something *unexpectedly* good happens, reward the features that were on  $\Rightarrow w_i \uparrow$ .
- Formal explanation: This is *online least squares with gradient descent*.
  - ▶ If **error** is Gaussian, then least squares = maximum likelihood.



# Example: Q-learning in Pacman

Credit: Dan Klein, Pieter Abbeel

- Initial value:  $Q(x, a) = 4.0f_{DOT}(x, a) - 1.0f_{GST}(x, a)$
- Features:  $f_i(x, a) = \frac{1}{dist\_pacman\_i}$        $\alpha = 0.004$



$$f_{DOT}(x, N) = 0.5 \quad f_{GST}(x, N) = 1$$
$$Q(x_t, N) = 1$$



# Example: Q-learning in Pacman

Credit: Dan Klein, Pieter Abbeel

- Initial value:  $Q(x, a) = 4.0f_{DOT}(x, a) - 1.0f_{GST}(x, a)$
- Features:  $f_i(x, a) = \frac{1}{dist\_pacman\_i}$        $\alpha = 0.004$



$R = -500$




$$f_{DOT}(x, N) = 0.5 \quad f_{GST}(x, N) = 1$$
$$Q(x_t, N) = 1$$

$$Q(x_{t+1}, \cdot) = 0$$



# Example: Q-learning in Pacman

Credit: Dan Klein, Pieter Abbeel

- Initial value:  $Q(x, a) = 4.0f_{DOT}(x, a) - 1.0f_{GST}(x, a)$
- Features:  $f_i(x, a) = \frac{1}{dist\_pacman\_i}$        $\alpha = 0.004$



$R = -500$



$$f_{DOT}(x, N) = 0.5 \quad f_{GST}(x, N) = 1 \\ Q(x_t, N) = 1$$

$$Q(x_{t+1}, \cdot) = 0$$

$$error = (R + \gamma \max_{a'} Q(x_{t+1}, a')) - Q(x_t, a_t) = -501$$

$$w_{DOT} \leftarrow 4.0 + \alpha(-501)0.5 \quad w_{GST} \leftarrow -1.0 + \alpha(-501)1.0$$

- Final value:  $Q(x, a) = 3.0f_{DOT}(x, a) - 3.0f_{GST}(x, a)$


# Online least squares

- Remember gradient descent for logistic regression (see Fundamentals).
- Using *online least squares* we can minimize the error, one point  $(x, y)_p$  at a time, by following the gradient  $w_m = w_m + \frac{1}{2}\alpha \frac{\partial J_p(w)}{\partial w_m}$ .
  - ▶ We have the same problems here: learning rate  $\alpha$ , overfitting, etc.

$$J_p(w) = \left( y - \sum_k w_k f_k(x) \right)^2$$

$$\frac{\partial J_p(x)}{\partial w_m} = -2 \left( y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \left( y - \sum_k w_k f_k(x) \right) f_m(x)$$



Credit: Alexander Amini, Daniela Rus

# Bibliography

- Richard S. Sutton, Andrew G. Barto. Reinforcement learning: An Introduction (second edition), 2018  
<http://incompleteideas.net/book/the-book.html>
- David Silver. Advanced Topics in Machine Learning, 2015  
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>
- Dan Klein, Pieter Abbeel. Artificial Intelligence CS188  
<http://ai.berkeley.edu>