



Reinforcement Learning

Machine Learning (69152)

Rubén Martínez Cantín

Dpto. Informática e Ingeniería de Sistemas
Universidad de Zaragoza

1 Introduction

- Markov Decision Processes (recap)
- Learning by interaction


2 Prediction

- Monte Carlo Evaluation
- Temporal Differences

3 Control

- Sampling control
- Exploration vs Exploitation
- Q-learning
- Q-learning with features

Markov Decision Processes




In particular, for **Markov Decision Processes (MDPs)**:

- Markov property $x_{t+1} = f(x_t)$
- Full observability $y_t = x_{t+1}$
- Action $a_t \leftarrow \pi(x_t)$

Markov Decision Process

- A Markov Decision Process is a tuple $\langle \mathcal{X}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$

- ▶ A set of states $x \in \mathcal{X}$.
- ▶ A set of actions $a \in \mathcal{A}$.
- ▶ A transition function
$$T(x, a, x') = p(x'|x, a)$$
- ▶ A reward function $R(x, a, x')$
- ▶ Maybe a start state and a terminal state.



Credit: Dan Klein, Pieter Abbeel

Markov Decision Process

- A Markov Decision Process is a tuple $\langle \mathcal{X}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$

- ▶ A set of states $x \in \mathcal{X}$.
- ▶ A set of actions $a \in \mathcal{A}$.
- ▶ A transition function
$$T(x, a, x') = p(x'|x, a)$$
- ▶ A reward function $R(x, a, x')$
- ▶ Maybe a start state and a terminal state.



- **Plot twist:** Now we don't know the transitions $T(\cdot)$ or rewards $R(\cdot)$
 - ▶ The agent must try actions to see the outcome.

Credit: Dan Klein, Pieter Abbeel

Learning by interaction

Model-based reinforcement learning

- Learn an approximate model based on experiences
- Assume that the learned model is correct and solve it as if known MDP.

Model-free reinforcement learning

- Learn the policy and values directly from interaction
- No explicit model of the transitions $T(\cdot)$ or rewards $R(\cdot)$

Example: Compute the expected age

We want to compute the expected age of all the students in the University.

Known Model $p(a)$

$$\mathbb{E}(A) = \sum_a p(a) * a$$

If we $p(a)$ is unknown we can ask several students for the age $[a_1, a_2, \dots]$

Model based

$$\hat{p}(a) = \frac{\text{num}(a)}{N}$$
$$\mathbb{E}(A) \approx \sum_a \hat{p}(a) * a$$

Model free

$$\mathbb{E}(A) \approx \frac{\sum_i a_i}{N}$$

Model based learning


- Model-Based Idea:
 - ▶ Learn an approximate model based on experiences
 - ▶ Solve for values as if the learned model were correct
- Step 1: Learn empirical MDP model
 - ▶ Count outcomes x' for each x, a
 - ▶ Normalize to give an estimate of $\hat{p}(x'|x, a)$
 - ▶ Discover each $\hat{R}(x, a, x')$ when we experience (x, a, x')
- Step 2: Solve the learned MDP
 - ▶ For example, use value iteration, as before



Credit: Dan Klein, Pieter Abbeel

Model-free learning

- **Model-free Idea:** why bother learning the transition or reward model?
 - ▶ Can we compute V , Q or π^* without a model?
 - ▶ **Monte Carlo:** Sample episodes, average rewards at the end.
 - ▶ **Temporal differences:** Use sampling to approximate the Bellman updates, compute new values during each learning step.



Credit: Dan Klein, Pieter Abbeel

Model-free learning

- Passive reinforcement learning
 - ▶ Evaluation given a policy π , find the value V^π .
 - ▶ *Learner* has no choice. Just follow the policy
 - ▶ The agent is doing **prediction**.
- Active reinforcement learning
 - ▶ Find the optimal policy/values: V^*, Q^*, π^* .
 - ▶ *Learner* can choose. Exploration vs Exploitation.
 - ▶ The agent is doing **control**.

Model-free learning

- Passive reinforcement learning
 - ▶ Evaluation given a policy π , find the value V^π .
 - ▶ *Learner* has no choice. Just follow the policy
 - ▶ The agent is doing **prediction**.
- Active reinforcement learning
 - ▶ Find the optimal policy/values: V^*, Q^*, π^* .
 - ▶ *Learner* can choose. Exploration vs Exploitation.
 - ▶ The agent is doing **control**.
- Remember: even if the learner can choose or not, it must take actions and interact with the world.

Monte Carlo Direct Evaluation

- Monte Carlo (MC) passive reinforcement learning.
- Goal: learn V^π from episodes of experience under policy π
- Act according to π .

$$x_1, a_1, R_2, x_2, \dots, x_n \sim \pi$$

- Recall that the utility is the total discounted reward:

$$U_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots R_n$$

- Recall that the value function is the expected return:

$$V^\pi(x) = \mathbb{E}_\pi[U_t | x_t = x]$$

- MC policy evaluation uses empirical mean return instead of expected return


$$V^\pi(x) = \sum_i U_t^{(i)} \mathbb{1}_{x_t^{(i)}=x} \quad \forall i \in \text{episodes}$$

Monte Carlo Direct Evaluation

- To evaluate state x
- Every time-step t that state x is visited in an episode,
- Increment counter $N(x) \leftarrow N(x) + 1$
- Increment total return $S(x) \leftarrow S(x) + U_t$
- Value is estimated by mean return $\hat{V}(x) = S(x)/N(x)$
- By law of large numbers, $\hat{V}(x) \rightarrow V^\pi(x)$ as $N(x) \rightarrow \infty$

Monte Carlo Direct Evaluation

Input policy π



Assume: $\gamma = 1$

Observed Episodes

Episode 1

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 3

- E, north, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 2

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 4

- E, north, C, -1
- C, east, A, -1
- A, exit, x, -10


Estimated values

	-10	
+8	+4	+10
	-2	

Credit: Dan Klein, Pieter Abbeel

Monte Carlo Direct Evaluation

Input policy π



Assume: $\gamma = 1$

Observed Episodes

Episode 1

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 3

- E, north, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 2

- B, east, C, -1
- C, east, D, -1
- D, exit, x, +10

Episode 4

- E, north, C, -1
- C, east, A, -1
- A, exit, x, -10

Estimated values

	-10	
+8 B	+4 C	+10 D
	-2 E	

- How can B and E have different values if both go to C for this policy?

Credit: Dan Klein, Pieter Abbeel

Incremental Monte Carlo

- Do we need to wait until the end of all episodes to compute the empirical mean?
- Remember the *expected age* example:
 - ▶ Model free approach: $\mathbb{E}(A) \approx \hat{A}_N = \frac{\sum_{i=1}^N a_i}{N}$
 - ▶ Incremental approach:

$$\begin{aligned}\hat{A}_N &\approx \frac{1}{N} \sum_{i=1}^N a_i \\&= \frac{1}{N} \left(a_N + \sum_{i=1}^{N-1} a_i \right) \\&= \frac{1}{N} \left(a_N + (N-1)\hat{A}_{N-1} \right) \\&= \hat{A}_{N-1} + \frac{1}{N} \left(a_N - \hat{A}_{N-1} \right)\end{aligned}$$

Incremental Monte Carlo for value function

- Update $V(x)$ incrementally after episode $x_1, a_1, R_2, x_2, \dots, x_n$
- For each state x_t with utility U_t

$$N(x_t) \leftarrow N(x_t) + 1$$

$$V(x_t) \leftarrow V(x_t) + \frac{1}{N(x_t)}(U_t - V(x_t))$$

- In non-stationary problems, it can be useful to track a running mean, that is, forget old episodes.

$$V(x_t) \leftarrow V(x_t) + \alpha(U_t - V(x_t))$$

Problems with Monte Carlo Direct Evaluation

- Samples are based on full episodes.
- The problem needs to be episodic. Not for continuous.
- It ignores MDP structure.
- Is there a way to exploit MDP structure and use Bellman equations?

Monte Carlo Policy Evaluation

- We had the Bellman expectation equation to update the value function:

$$V^{\pi}(x) = \sum_{a \in \mathcal{A}} \pi(a|x) \left(\sum_{x' \in \mathcal{X}} p(x'|x, a) (R(x, a, x') + \gamma V^{\pi}(x')) \right)$$

- This approach exploits the MDP structure and is valid for continuous problems. 😊
- But it requires $p(x'|x, a)$ and $R(x, a, x')$ to do it. 😞

Sample based policy evaluation

- Instead of sampling full episodes, we can sample the value update.

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(a|x) \left(\sum_{x' \in \mathcal{X}} p(x'|x, a) (R(x, a, x') + \gamma V^\pi(x')) \right)$$

- From one state x , we sample just the next action $a_1, a_2, \dots \sim \pi(x)$ and the next transition $x_1 \sim p(x'|a_1, x), x_2 \sim p(x'|a_2, x), \dots$
- Then, we collect the outcome for the action/transition:

$$sample_1 = R(x, a_1, x'_1) + \gamma V_k^\pi(x'_1)$$

$$sample_2 = R(x, a_2, x'_2) + \gamma V_k^\pi(x'_2)$$

...

- Finally,

$$V_{k+1}^\pi(x) \leftarrow \frac{1}{N} \sum_{i=1}^N sample_i$$

Temporal Difference

- **Idea:** learn from any interaction.
- Combine the sample based policy evaluation with the incremental Monte Carlo method.
- Remember: Incremental Monte Carlo updates the value with the *actual* utility U_t .

$$V(x_t) \leftarrow V(x_t) + \alpha(U_t - V(x_t))$$

- Temporal Difference updates the value with the *estimated* utility $R_{t+1} + \gamma V(x_{t+1})$

$$V(x_t) \leftarrow V(x_t) + \alpha(R_{t+1} + \gamma V(x_{t+1}) - V(x_t))$$

Temporal Difference

- Temporal Difference (TD) can be learned after any interaction. No need to wait for the end of the episode.
- It can be used for continuous problems.
- It can even learn from incomplete sequences.
- It exploits MDP structure.

$$\begin{aligned} \text{sample}_k &= R(x, a, x') + \gamma V_k^\pi(x') \\ V_{k+1}^\pi(x) &\leftarrow (1 - \alpha) V_k^\pi(x) + \alpha \cdot \text{sample}_k = \\ &= V_k^\pi(x) + \alpha(\text{sample}_k - V_k^\pi(x)) \end{aligned}$$

- Remember that $\alpha \in [0, 1]$ is used to *forget*.
 - $\gamma \rightarrow 0$ ignores the distant future.
 - $\alpha \rightarrow 1$ forgets the distant past

Forgetting the past

- Remember, we can use a running mean to forget old updates.
 - In MC, for non-stationary problems
 - In TD, recursive update of Bellman equation:

$$V(x_t) \leftarrow V(x_t) + \alpha(\text{update}_t - V(x_t))$$

- A running average $\hat{z}$ of $\{z\}_{i=1}^N$ elements:

$$\hat{z}_k = \hat{z}_{k-1} + \alpha(z_k - \hat{z}_{k-1}) = (1 - \alpha)\hat{z}_{k-1} + \alpha z_k$$

- Weights recent samples more:


$$\hat{z}_k = \frac{z_k + (1 - \alpha)z_{k-1} + (1 - \alpha)^2 z_{k-2} + \dots}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots}$$

MC vs TD

- Both methods converge to the true value if experience $\rightarrow \infty$
- MC requires full episodes.
- MC is a very simple and general idea. It can be applied for any system.
- TD exploits the Markov property.
 - ▶ It is more efficient for Markov environments.
 - ▶ It might not converge for certain environments (e.g.: non-Markov).

Monte Carlo Update


$$V(x_t) \leftarrow V(x_t) + \alpha(U_t - V(x_t))$$



Credit: David Silver

Temporal Difference Update


$$V(x_t) \leftarrow V(x_t) + \alpha(R_{t+1} + \gamma V(x_{t+1}) - V(x_t))$$



Credit: David Silver

Dynamic Programming Update


$$V(x_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(x_{t+1})]$$



Credit: David Silver

Unified view of RL

- The TD algorithm that we have seen is called $TD(0)$, because it considers only one step ahead.
- The generalization is called $TD(\lambda)$, which combines multiple steps ahead.



Credit: David Silver

Sampling control

- We have seen how to use MC and TD for **prediction** (policy evaluation).
- We are going to use MC and TD for **control** (optimal policy/value).
- **On-policy** learning
 - ▶ *Learn on the job*
 - ▶ Learn about policy π from experience sampled from π
- **Off-policy** learning
 - ▶ *Look over someone's shoulder*
 - ▶ Learn about policy π from experience sampled from μ

Quick Parenthesis: Q-values

- Computing the greedy policy from $V(x)$ requires the MDP model $p(x'|x, a)$ and $R(x, a, x')$.

$$\pi'(x) = \arg \max_a \sum_{x'} p(x'|x, a) (R(x, a, x') + \gamma V(x'))$$



- Computing the greedy policy from $Q(x, a)$ can be done model-free:

$$\pi'(x) = \arg \max_a Q(x, a)$$

- **Solution:** we are going to do control based on Q-functions.


SARSA

- Iteratively estimate Q and π (similar to policy iteration)
 - Policy evaluation (estimate Q given π):
 - Good: Monte Carlo: run multiple full episodes, then update policy.
 - Better: Improve policy **after each episode**.
 - Best: Do not run full episodes. Use Temporal Differences.
- $$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R + \gamma Q(x_{t+1}, a_{t+1}) - Q(x_t, a_t))$$
- Policy improvement (improve π given Q):
 - Exploit: The optimal policy is **greedy** to the value function...
 - Explore: ... but remember that we also need to **explore**!



Exploration vs Exploitation

- The greedy policy is the best option if we know the model.
- Now, we are learning by experience.
 - ▶ The agent must have diverse experiences to learn.
 - ▶ But the agent must also find an optimal behavior.



Regret

- Regret measures your **mistakes**. The (expected) reward of your actions, including suboptimal choices, against the (expected) optimal reward.
- Minimizing regret is not only learning to be optimal. It is **optimally learning to be optimal**.
- Pure random exploration will always find the optimal policy, but it has very high regret.
- A good trade-off of exploration and exploitation can provide **optimal regret**.






Exploration-exploitation trade-off principles

- The limit cases are the *greedy policy* (always exploit) and the *random policy* (always exploit).
 - ▶ We have seen that both policies have high regret.
- Random mixing (ϵ -greedy)
 - ▶ The simplest approach: flip a coin and choose *greedy* or *random policy* depending on the outcome.
- Optimistic initialization
 - ▶ Assume *unknown* = *best*. Try everything at least once.
- Optimism in the Face of Uncertainty
 - ▶ In case of uncertainty, assume the best possible outcome.
 - ▶ For example, upper confidence bound
 - ▶ If you are right: you win!
 - ▶ If you are wrong: you learn a lot!

Exploration-exploitation demo

Exercise/demo: <https://drive.google.com/file/d/1xAph2c1P8pPcJudiWBCd9Rrqt4p2E0Ra/view?usp=sharing>



DP vs TD algorithms

	Full backup	TD Backup
$V^\pi(x), Q^\pi(x, a)$	Policy evaluation	TD learning
$\pi^*(x, a), V^*(x), Q^*(x, a)$	Policy iteration	SARSA
$V^*(x), Q^*(x, a)$	Value iteration	???

DP vs TD algorithms

	Full backup	TD Backup
$V^\pi(x), Q^\pi(x, a)$	Policy evaluation	TD learning
$\pi^*(x, a), V^*(x), Q^*(x, a)$	Policy iteration	SARSA
$V^*(x), Q^*(x, a)$	Value iteration	Q-learning

Q-learning

- Start with Q-value iteration

$$Q_{k+1}(x_t, a_t) = \sum_{x_{t+1}} p(x_{t+1}|x_t, a_t) \left(R(x_t, a_t, x_{t+1}) + \gamma \max_{a'} Q_k(x_{t+1}, a') \right)$$

- Consider the sample update

$$sample = R(x_t, a_t, x_{t+1}) + \gamma \max_{a'} Q_k(x_{t+1}, a')$$

- Which, using the running average, can be incorporated as:

$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha (R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t))$$

Q-learning algorithm

Algorithm 1 Q-learning algorithm

Input: Step size $\alpha \in [0, 1]$, policy parameters: e.g., small $\epsilon > 0$.

Initialize $Q(x, a)$, for all $x \in \mathcal{X}$, $a \in \mathcal{A}$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

for each episode **do**:

 Initialize x_0

for each step t in episode **do**:

 Choose a_t from x_t using a policy derived from Q (e.g.: ϵ -greedy)

 Take action a_t , observe R_{t+1} and x_{t+1} .

$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t))$

if x_{t+1} is terminal **then** Stop episode

On-policy vs Off-policy

- SARSA and Q-learning require 2 actions for each update.

$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha(R_{t+1} + \gamma Q(x_{t+1}, a') - Q(x_t, a_t))$$

- This may come from different sources:
 - ▶ The **behavior policy** is how the agent is acting.
 - ▶ The **target policy** is the policy that the agent is learning.
- On-policy vs off-policy
 - ▶ On-policy methods use the same policy for the behavior and target.
 - ▶ Off-policy methods use a different policy. They can learn the optimal policy even acting suboptimally!

On-policy vs Off-policy

- In SARSA, all actions are selected according to the target policy π , for example: ϵ -greedy.


$$Q(x_t, \mathbf{a}_t) \leftarrow Q(x_t, \mathbf{a}_t) + \alpha(R_{t+1} + \gamma Q(x_{t+1}, \mathbf{a}_{t+1}) - Q(x_t, \mathbf{a}_t))$$

- In Q-learning, actions are selected according to the **behavior policy** μ , for example: ϵ -greedy, but the agent assumes that in the future it will be optimal (π is the **greedy policy**).

$$Q(x_t, \mathbf{a}_t) \leftarrow Q(x_t, \mathbf{a}_t) + \alpha(R_{t+1} + \gamma \max_{\mathbf{a}'} Q(x_{t+1}, \mathbf{a}') - Q(x_t, \mathbf{a}_t))$$


Scaling Q-learning

- How many states does smallGrid Pacman have?



Scaling Q-learning

- How many states does smallGrid Pacman have?



- Classic Pacman has: $240^5 \cdot 2^{240} = 1.766847110^{72}$ states
- Basic Q-learning keeps a table with a Q-value for every state and action.
- This idea is not able to scale:
 - ▶ During training, it is impossible to visit all the states/actions.
 - ▶ The table cannot be kept in memory.

Generalization Q-learning

- We want to learn from few states and generalize to **similar states**.


We learn that this is a
bad state



Basic Q-learning tell
us nothing about this
state.



Or even this state!



- What do they have in common?

Credit: Dan Klein, Pieter Abbeel

Approximate Q-learning

- Idea: exploit certain properties (features) of state.
- A feature is a function from states or (state,action) to numbers.
- We can use features to describe other functions
- Examples of features:
 - ▶ Distance to a ghost.
 - ▶ Distance to a dot.
 - ▶ Number of ghost.
 - ▶ Is Pacman in a corner/tunnel?
 - ▶ $1/(distance_to_dot)^2$
 - ▶ ...
 - ▶ Is it the state in this slide?
 - ▶ ...
 - ▶ Does this action get me closer to food?



Credit: Dan Klein, Pieter Abbeel

Linear function approximation

- Using a feature representation, we can write a q function (or value function) for any state using a few weights:

$$V(x) = w_1 f_1(x) + w_2 f_2(x) + \dots + w_n f_n(x)$$

$$Q(x, a) = w'_1 f_1(x, a) + w'_2 f_2(x, a) + \dots + w'_n f_n(x, a)$$

- Advantage: our experience is summed up in a few powerful numbers
- Disadvantage: states may share features but actually be very different in value!

Approximate Q-learning

- Take a linear Q-value function:

$$Q(x, a) = w_1 f_1(x, a) + w_2 f_2(x, a) + \dots + w_n f_n(x, a)$$

- Q-learning update is based on estimation error:

$$\text{error} = \left(R_{t+1} + \gamma \max_{a'} Q(x_{t+1}, a') \right) - Q(x, a)$$

- Basic Q-learning updates each value with the error:

$$Q(x, a) \leftarrow Q(x, a) + \alpha \cdot \text{error}$$

- Approximate Q-learning updates the weights of active features:

$$w_i \leftarrow w_i + \alpha \cdot \text{error} \cdot f_i(x, a)$$

Approximate Q-learning

- Approximate Q-learning updates the weights of active features:

$$w_i \leftarrow w_i + \alpha \cdot \text{error} \cdot f_i(x, a)$$

- Intuitive interpretation:
 - ▶ If something *unexpectedly* bad happens, penalize the features that were on $\Rightarrow w_i \downarrow$.
 - ▶ If something *unexpectedly* good happens, reward the features that were on $\Rightarrow w_i \uparrow$.
- Formal explanation: This is *online least squares with gradient descent*.
 - ▶ If **error** is Gaussian, then least squares = maximum likelihood.

Example: Q-learning in Pacman

Credit: Dan Klein, Pieter Abbeel

- Initial value: $Q(x, a) = 4.0f_{DOT}(x, a) - 1.0f_{GST}(x, a)$
- Features: $f_i(x, a) = \frac{1}{dist_pacman_i}$ $\alpha = 0.004$



$$f_{DOT}(x, N) = 0.5 \quad f_{GST}(x, N) = 1$$

$$Q(x_t, N) = 1$$

Example: Q-learning in Pacman



Credit: Dan Klein, Pieter Abbeel

- Initial value: $Q(x, a) = 4.0f_{DOT}(x, a) - 1.0f_{GST}(x, a)$
- Features: $f_i(x, a) = \frac{1}{dist_pacman_i}$ $\alpha = 0.004$



$$f_{DOT}(x, N) = 0.5 \quad f_{GST}(x, N) = 1$$
$$Q(x_t, N) = 1$$

$$R = -500$$



$$Q(x_{t+1}, \cdot) = 0$$


Example: Q-learning in Pacman

Credit: Dan Klein, Pieter Abbeel


- Initial value: $Q(x, a) = 4.0f_{DOT}(x, a) - 1.0f_{GST}(x, a)$
- Features: $f_i(x, a) = \frac{1}{\text{dist_pacman_i}}$ $\alpha = 0.004$



$R = -500$



$a_t = \text{NORTH}$



$$f_{DOT}(x, N) = 0.5 \quad f_{GST}(x, N) = 1 \\ Q(x_t, N) = 1$$

$$Q(x_{t+1}, \cdot) = 0$$

$$\text{error} = (R + \gamma \max_{a'} Q(x_{t+1}, a')) - Q(x_t, a_t) = -501$$

$$w_{DOT} \leftarrow 4.0 + \alpha(-501)0.5 \quad w_{GST} \leftarrow -1.0 + \alpha(-501)1.0$$

- Final value: $Q(x, a) = 3.0f_{DOT}(x, a) - 3.0f_{GST}(x, a)$


Online least squares

- Remember gradient descent for logistic regression (see Fundamentals).
- Using *online least squares* we can minimize the error, one point $(x, y)_p$ at a time, by following the gradient $w_m = w_m + \frac{1}{2}\alpha \frac{\partial J_p(w)}{\partial w_m}$.
 - We have the same problems here: learning rate α , overfitting, etc.

$$J_p(w) = \left(y - \sum_k w_k f_k(x) \right)^2$$

$$\frac{\partial J_p(x)}{\partial w_m} = -2 \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$



Credit: Alexander Amini, Daniela Rus

Bibliography

- Richard S. Sutton, Andrew G. Barto. Reinforcement learning: An Introduction (second edition), 2018
<http://incompleteideas.net/book/the-book.html>
- David Silver. Advanced Topics in Machine Learning, 2015
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>
- Dan Klein, Pieter Abbeel. Artificial Intelligence CS188
<http://ai.berkeley.edu>