

Machine Learning - Deep Learning fundamentals (69152) CNNs

Master in Robotics, Graphics and Computer Vision
Ana C. Murillo



Universidad
Zaragoza


Today

- What's a CNN?
- Well-known CNN architectures

CNNs

- Convolutional (deep) Neural Networks (CNN)

[DEMO online, CNN](#)




**NOT all LAYERS
are the same**


CNNs - Basic Layers

- **Input:** raw pixel values. RGB image 32x32, volume [32x32x3]
- **Convolutional:** compute output of neurons connected to local input regions (each computes dot product between their weights and the region). Larger volume [32x32x12].
- **RELU:** element-wise activation function (e.g. thresholding at zero). Volume unchanged ([32x32x12]).
- **Pooling:** downsampling operation. e.g. reduce volume to [16x16x12].
- **Fully-connected:** compute class scores. Each neuron in this layer will be connected to all the numbers in the previous volume. Volume of size [1x1x10].

CNNs - pooling example



CNNs - pooling example




CNNs

- A bit of “history” about CNNs

CNNs - Image Classification

CNNs *basics*: AlexNet (Beginning of the *new CNN boom*):
a layered model composed of convolution, subsampling, and further operations followed by a holistic representation.



***Illustrated the
benefits of CNNs***

- 15 million annotated images
- over 22000 classes
- **ReLU, data augmentation, dropout**


IMAGENET

CNNs

CNNs *basics*: DeConvNet




image patches that strongly activate 1st layer filters



1st layer filters

***better understanding:
earlier vs later layers***


- Similar architecture to AlexNet
- less training data - smaller filters
- Deconvolutional Network - **visualization**



conv₅ DeConv visualization
[Zeiler-Fergus]

CNNs

CNNs *basics*: DeConvNet




1st layer filters



CNNs

CNNs *basics*: DeConvNet

top 9 activations for a few feature maps
(projection to pixel space using DeConv
and actual image patches)




conv₅ DeConv visualization
[Zeiler-Fergus]

CNNs - Image Classification + Detection

- **Analyze the *feature maps*: Region - CNN**


(R-CNN - 2013, Fast R-CNN - 2015, Faster R-CNN - 2016)



**Analyze feature maps
(activations) to learn
where the objects are**

CNNs - Image classification

- *Deeper* classification: VGG - **very deep** ConvNets
 - Smaller filters - less params
 - more depth! —> always good (?)
 - Consecutive conv. layers (smaller filters).
 - Compare to AlexNet 5 conv. layers



How deep?





Image classification (deeper)



16 convolutional layers



50 convolutional layers



152 convolutional layers



Krizhevsky, I. Sutskever, and G. E. Hinton. *ImageNet classification with deep convolutional neural networks*. In Proc. NIPS, 2012.

C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabbinovich. *Going deeper with convolutions*. In Proc. CVPR, 2015.

K. Simonyan and A. Zisserman. *Very deep convolutional networks for large-scale image recognition*. In Proc. ICLR, 2015.

K. He, X. Zhang, S. Ren, and J. Sun. *Deep residual learning for image recognition*. In Proc. CVPR, 2016.

Are they just deeper version of the same?

Deeper: How does memory get affected?

INPUT: [224x224x3] memory: $224*224*3=150K$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800K$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400K$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200K$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100K$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25K$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$


FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$

TOTAL memory: $24M * 4 \text{ bytes} \approx 96MB$ / image (only forward! $\sim *2$ for bwd)

TOTAL params: 138M parameters

Very Deep Convolutional Networks for Large-Scale Image Recognition.
K. Simonyan, A. Zisserman. 2014



Deeper: How does memory get affected?

INPUT: [224x224x3] memory: 224*224*3=150K params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: **224*224*64=3.2M** params: (3*3*3)*64 = 1,728

CONV3-64: [224x224x64] memory: **224*224*64=3.2M** params: (3*3*64)*64 = 36,864

POOL2: [112x112x64] memory: 112*112*64=800K params: 0

CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*64)*128 = 73,728

CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*128)*128 = 147,456

POOL2: [56x56x128] memory: 56*56*128=400K params: 0

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*128)*256 = 294,912

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824

POOL2: [28x28x256] memory: 28*28*256=200K params: 0

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*256)*512 = 1,179,648

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2,359,296

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2,359,296

POOL2: [14x14x512] memory: 14*14*512=100K params: 0

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

POOL2: [7x7x512] memory: 7*7*512=25K params: 0

FC: [1x1x4096] memory: 4096 params: 7*7*512*4096 = **102,760,448**

FC: [1x1x4096] memory: 4096 params: 4096*4096 = 16,777,216

FC: [1x1x1000] memory: 1000 params: 4096*1000 = 4,096,000

Note:

Most memory is in early CONV

Most params are in late FC

image

Conv-64

Conv-64

maxpool

Conv-128

Conv-128

maxpool

Conv-256

Conv-256

maxpool

Conv-512

Conv-512

Conv-512

maxpool

Conv-512

Conv-512

Conv-512

maxpool

fc-4096

fc-4096

fc-2622


Softmax

slide from
Fei-Fei Li & Justin Johnson & Serena Yeung
<http://cs231n.stanford.edu>


**What is more relevant for
batch size design
decisions?**

Image classification (deeper)

- Not only more layers! **it's getting hard(er) to train ...**



ImageNet Classification with Deep Convolutional Neural Networks
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012



16 convolutional layers

50 convolutional layers

152 convolutional layers

Krizhevsky, I. Sutskever, and G. E. Hinton. *ImageNet classification with deep convolutional neural networks*. In Proc. NIPS, 2012.


C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich. *Going deeper with convolutions*. In Proc. CVPR, 2015.

K. Simonyan and A. Zisserman. *Very deep convolutional networks for large-scale image recognition*. In Proc. ICLR, 2015.

K. He, X. Zhang, S. Ren, and J. Sun. *Deep residual learning for image recognition*. In Proc. CVPR, 2016.


Some *reference modules* for CNN architectures ...

- to train better: GoogLeNet, ResNet, DenseNet, ...




ImageNet Classification with Deep Convolutional Neural Networks
Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton. NIPS 2012

Inception module



Going Deeper With Convolutions.
C.Szegedy, et al. CVPR 2015.

Residual module




Deep residual learning for image recognition.
He, Kaiming, et al. CVPR 2016



Not just
“linear” list of
layers anymore

Dense module



Densely Connected Convolutional Networks.
Huang, G., et al. CVPR 2017.

Open or interesting problems related to deep learning?

Not enough data to train? (or resources)


CNNs & Transfer Learning

- CNNs are able to generalize well!

- great **features**

- **fine-tuning**

Lab 2
you'll practice
some of this




deep features

Transfer Learning


- What is it?
- When would you use transfer learning?
- Types of transfer learning?

Transfer Learning: features

- E.g., features from AlexNet last layer: 4096 dims. vector




Test image L2 Nearest neighbors in feature space



How do we get these features?


Transfer Learning: features

- E.g., features from AlexNet last layer: 4096 dims. vector




How do we use these features?

Test image L2 Nearest neighbors in feature space



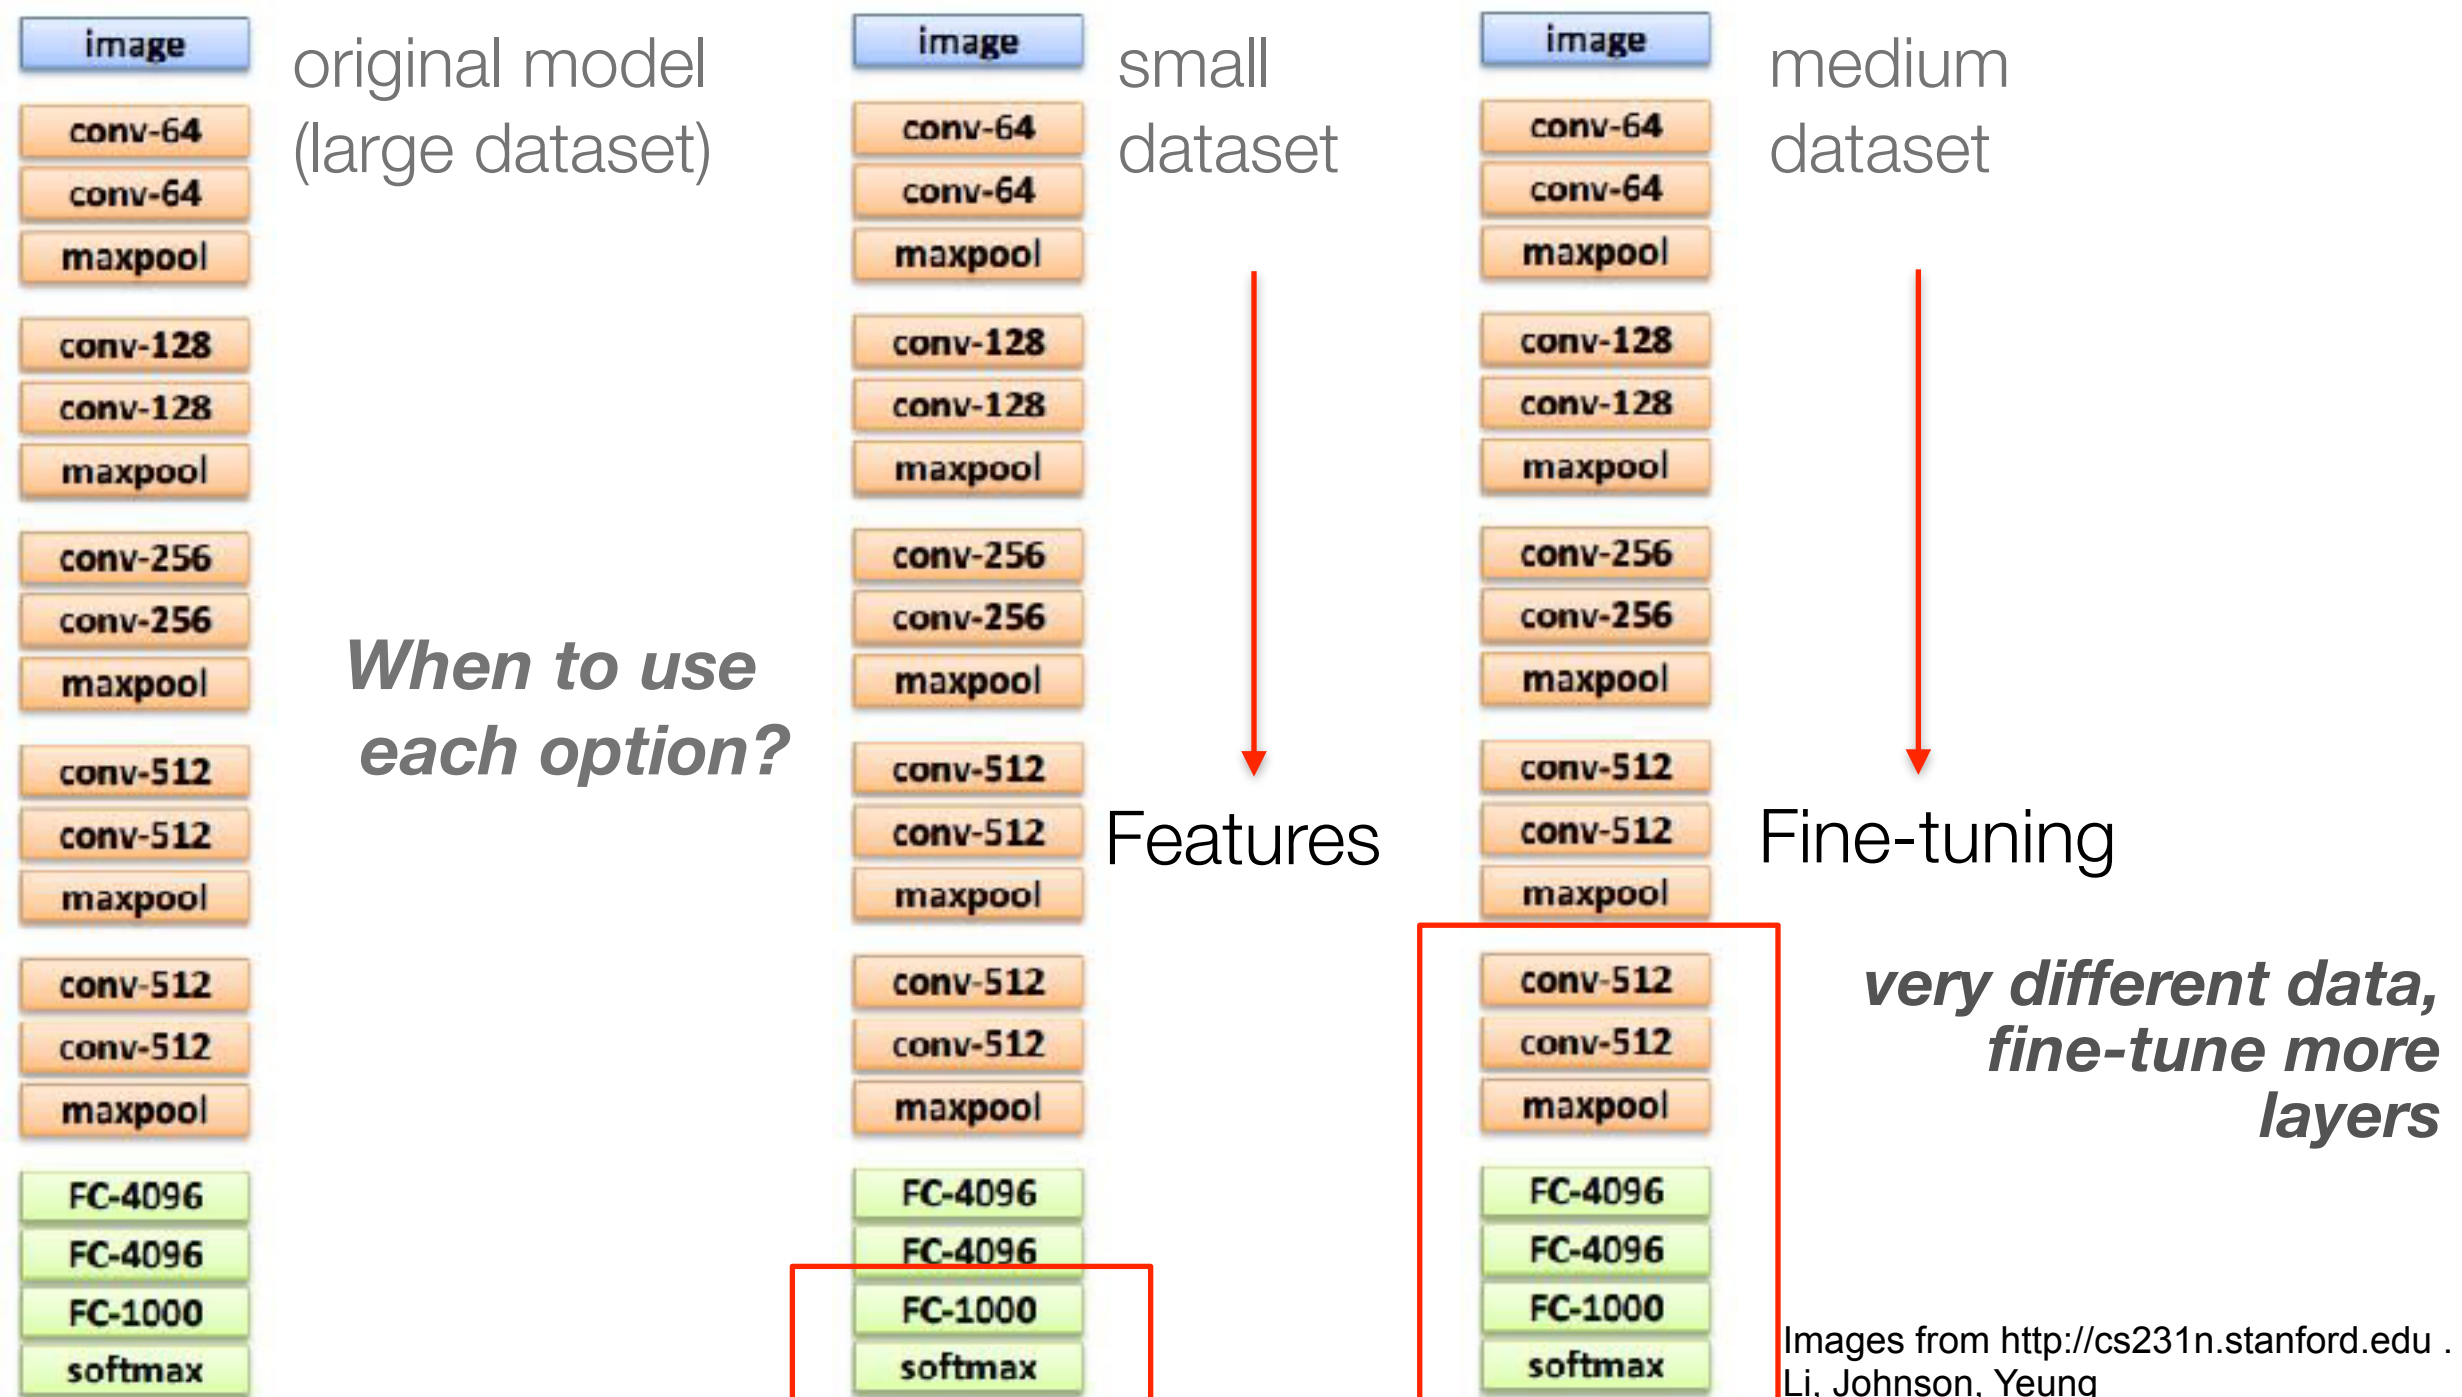
Transfer Learning: fine-tune

- For example, fine-tune ImageNet AlexNet for non Image-Net classes
- Basically, initialise weights to something more “*interesting*” than random (careful also with hyperparameters!)



Transfer Learning

... when not enough resources to train from scratch



Transfer Learning: fine-tuning

- not just that! very very spread out strategy:

Object Detection (Fast R-CNN)





Image captioning: CNN + RNN



*Pre-trained
models to
initialize certain
blocks*

Karpathy and Fei-Fei, "Deep Visual-Semantic Alignments for Generating Image Descriptions", CVPR 2015
Figure copyright IEEE, 2015. Reproduced for educational purposes.

Girshick, "Fast R-CNN", ICCV 2015
Figure copyright Ross Girshick, 2015. Reproduced with permission.

Lab 2: NN and CNNs applied in Computer Vision

- Understand a **DNN implemented from scratch**
- Implement a toy CNN with Keras
- Fine-tune a well-known CNN architecture with Keras

All this is applied to Image Classification

- Optional tasks:
 - implement variations of the classification networks
 - explore object recognition model
 - explore semantic segmentation model

***Available in
Moodle.***

***Please check
the prior-work
task (prepare
your set up
and data)***

TO-DO ...

Have you all used COLAB?

- Lab 2 - **TOMORROW (18 OCT , 20 OCT) —> WED. @ A07 / FRIDAY @ L0.06a ADA BYRON**
- **PLEASE have your computer ready with Tensorflow2+Keras (ideally with GPU available) AND/OR we will use Google COLAB**
- Recommended COLAB tutorial if you have not used it much before:
<https://colab.research.google.com/notebooks/intro.ipynb>
https://colab.research.google.com/notebooks/basic_features_overview.ipynb
[Loading data: Drive, Sheets, and Google Cloud Storage](#)

ASSIGNMENT BEFORE YOUR LAB:

- [illegible]

```
data/  
  dogs/  
    dog001.jpg  
    dog002.jpg  
    ...  
  cats/  
    cat001.jpg  
    cat002.jpg  
    ...
```


Lab 2: NN and CNNs applied in Computer Vision

- Frameworks
 - Caffe, **Tensorflow+Keras**, **Pytorch**, ...
- Model **zoos**
 - in every framework (*you should always start by looking into existing models. “rare” to “require” to implement a network from scratch*)


Understanding the training process

Initial sanity checks (after a few “toy” iterations)


- *Correct* loss at chance
- Increasing regularisation strength increases loss
- Overfit a tiny training set (you can set regularisation strength to zero to better see this)

Understanding the training process

- Monitor the loss function




Typical loss function.
Batch size a little low? (noisy plot)





Understanding the training process

- Monitor train/val accuracy



Understanding the training process

- Visualize first conv layer weights



Understanding the training process

Additional tips:

- **Decay your learning rate** during training
- If you can afford it
 - **Search** for good **hyperparameters** with random search (not grid search). From coarse to fine
 - Consider **model ensembles** (not in our labs)

Lab 2: NN and CNNs applied in Computer Vision

- Understand a **DNN implemented from scratch**
- Implement a toy CNN with Keras
- Fine-tune a well-known CNN architecture with Keras

All this is applied to Image Classification

- Optional tasks:
 - implement variations of the classification networks
 - explore object recognition model
 - explore semantic segmentation model

***Exercise now
(start for Lab2)***

[COLAB](#)