

Machine Learning - Deep Learning fundamentals (69152) CNNs

Master in Robotics, Graphics and Computer Vision
Ana C. Murillo



Universidad
Zaragoza

Today

- Open problems with deep learning?
- More well-known architectures

Open or interesting problems related to deep learning?

More architectures ...


- Not only classification —> detection?

Deep Learning & Regression

- **Detection**

predict normalized b-box coordinates


$$p = (x_1, y_1, x_2, y_2)$$



Fast R-CNN. Ross Girshick. 2015

Different versions of Region-CNN:

- R-CNN - 2013
- Fast R-CNN - 2015
- Faster R-CNN - 2016



CLASSIFICATION + DETECTION

Analyze feature maps (activations) to learn where the objects are

Deep Learning & Regression


- Linear regression: relationship between a scalar (y) and one or more explanatory variables (x)

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip} + \varepsilon_i = \mathbf{x}_i^T \boldsymbol{\beta} + \varepsilon_i, \quad i=1, \dots, n$$

“map” to a continuous output (regression)

VS

“map” to a discrete output (classification)




Wikipedia:
File:Linear least squares example2.png

Deep Learning & Regression

- Regression of **Pose** - *PoseNet*
 - Train the network to output 3D position ($\mathbf{x}$) and orientation ($\mathbf{q}$) (7 dims). Loss function:

$$loss(I) = \|\hat{\mathbf{x}} - \mathbf{x}\|_2 + \beta \left\| \hat{\mathbf{q}} - \frac{\mathbf{q}}{\|\mathbf{q}\|} \right\|_2$$

- Modify GoogLeNet:




Deep Learning & Regression

- Regression of **Pose** - *PoseNet*

- Train the network to output 3D position ($\mathbf{x}$) and orientation ($\mathbf{q}$) (7 dims). Loss function:

$$loss(I) = \|\hat{\mathbf{x}} - \mathbf{x}\|_2 + \beta \left\| \hat{\mathbf{q}} - \frac{\mathbf{q}}{\|\mathbf{q}\|} \right\|_2$$

- Modify GoogLeNet. Key change: replace softmax classifiers with affine regressors (each final fully connected layer now outputs a pose vector of 7-dims).




More architectures ...

More accurate image understanding?


Deep Learning & pixel classification

- Dense labeling/**Segmentation** - Fully Convolutional Net (FCN)



Deep Learning & pixel classification

- Dense labeling/**Segmentation** - Encoder-Decoder




<http://mi.eng.cam.ac.uk/projects/segnet/>

Deep Learning & pixel classification

- **Detection + Instance Segmentation**

Mask R-CNN - 2017

(R-CNN + Semantic Segmentation)




More architectures ...

- More details? Interaction? Time?
- More “complex” input data


Deep Learning & Video

- Fuse multiple frame info
- high-resolution center crop + low-resolution full image (reduced input size, more efficient training)




Deep Learning & Video

- Fuse multiple frame info. Many strategies



Deep Learning & Multi-Modal

- Multi-modal input from images
 - Deep learning features: Pose-based CNN
 - Combine **parts**, **pose** and **flow** with CNNs (aggregates motion and appearance information)



Deep Learning & Video

Do you see any problem with this way of treating sequential/video data?

Demos to understand CNNs

- karpathy : [DEMO online, CNN](#)
- [Playground tensorflow](#)
- visualisation (places CNN)

<http://people.csail.mit.edu/torralba/research/drawCNN/drawNet.html>

Visualizing and Understanding Convolutional Neural Networks
Matthew Zeiler and Rob Fergus. **2013**

Later ...

- More advanced models
- ... and different supervision strategies
 - DRL
 - Unsupervised
 - Recurrent architectures

Bibliography - Resources for some of the materials today

- Stanford classes on deep learning for Computer Vision (<http://cs231n.stanford.edu>) and Deep Learning (<https://cs230.stanford.edu/>)
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, Deep Learning, MIT Press, 2016. <http://www.deeplearningbook.org>
- Deep Learning Summer School Montreal: <https://mila.quebec/en/cours/deep-learning-summer-school-2017/>